

# Statistics Templates

*Go Ito*

*August 9th, 2019*

## Contents

<b>Package Installation</b>	<b>3</b>
<b>Generic library to use</b>	<b>3</b>
<b>Useful functions</b>	<b>3</b>
<b>Data Reading</b>	<b>4</b>
Built in data . . . . .	4
.txt file . . . . .	4
.csv file . . . . .	4
.xls / .xlsx file . . . . .	5
.json file . . . . .	5
SQL . . . . .	5
<b>Exploratory Analysis</b>	<b>6</b>
Summary . . . . .	6
Histogram / Distributions . . . . .	7
Boxplots . . . . .	9
Plots of two variables . . . . .	10
Interaction Plots . . . . .	10
<b>Data Analysis, Prediction, and Classification</b>	<b>12</b>
Simple Linear Regression . . . . .	12
Assumption Check . . . . .	12
Statistical Inference . . . . .	13
ANOVA SLR . . . . .	14
Transformation . . . . .	16
Weighted Least Square . . . . .	18
Multiple Linear Regression . . . . .	19
Assumptions . . . . .	19
Statistical Inference . . . . .	21
Model Diagnosis : Added Variable Plot . . . . .	22
ANOVA MLR . . . . .	23
Nested Model and Partial F-test . . . . .	23
ANCOVA (Analysis of Covariance) . . . . .	24
Polynomial Regression . . . . .	26
Model Selection Criteria . . . . .	31
Subset Selection: Best Subset Model . . . . .	31
Subset Selection: Stepwise Regression . . . . .	35
Shrinkage: Regularization . . . . .	40
Dimension Reduction: Principle Component Analysis . . . . .	45
Spline . . . . .	48
Generalized Additive Models . . . . .	51
Logistic Regression . . . . .	53
Bayes Classifier . . . . .	57
K Nearest Neighbor . . . . .	58
Linear Discriminant Analysis . . . . .	61
Quadratic Discriminant Analysis . . . . .	67

Support Vector Machine . . . . .	68
K-Means Clustering . . . . .	69
Kernelized Clustering . . . . .	70
EM Type Algorithm . . . . .	71
Decision Trees . . . . .	72
Random Forest . . . . .	75
Neural Network . . . . .	80
Comparing all classification methods . . . . .	81
<b>Experimental Design</b>	<b>82</b>
t-test . . . . .	82
Completely Randomized Design . . . . .	83
Paired Comparison Design . . . . .	85
ANOVA for experiment . . . . .	87
Basics . . . . .	87
Model Adequacy . . . . .	89
Post-hoc analysis . . . . .	89
Power of the test . . . . .	89
Completely Randomized Block Design . . . . .	91
Latin Square Design . . . . .	92
Greco-Latin Square Design . . . . .	92
Balanced Incomplete Design . . . . .	94
Factorial Design with no blocking . . . . .	96
General Factorial Design with blocking . . . . .	99
$2^k$ Factorial Design . . . . .	102
$2^k$ Factorial Design with blockings (confoundings) . . . . .	105
$2^{k-p}$ Fractional Factorial Design . . . . .	107
<b>Computational Techniques</b>	<b>108</b>
Standarization / Normalization . . . . .	108
K-Fold Cross Validation . . . . .	109
Bootstrap . . . . .	111
<b>R Features</b>	<b>112</b>
Basics . . . . .	112
Functions . . . . .	112
I() and poly() for polynomial regression . . . . .	113
Normalization and Standarization . . . . .	115
apply(), sapply(), tapply(), lapply() . . . . .	115
Package: dplyr . . . . .	115

## Package Installation

```
#install.packages("ggplot2")
#install.packages("dplyr")
#install.packages("tidyverse")
#install.packages("readxl")
#install.packages("alr3")
#install.packages("MASS")
#install.packages("ISLR")
#install.packages("class")
#install.packages("caret")
#install.packages("e1071")
#install.packages("leaps")
#install.packages("boot")
#install.packages("crossval")
#install.packages("resample")
#install.packages("glmnet")
#install.packages("pls")
#install.packages("splines")
#install.packages("gam")
#install.packages("akima")
#install.packages("tree")
#install.packages("randomForest")
#install.packages("nnet")
#install.packages("NeuralNetTools")
#install.packages("RSQLite")
#install.packages("Hmisc")
#install.packages("pwr")
#install.packages("agricolae")
#install.packages("crossdes")
#install.packages("phia")
#install.packages("gplots")
#install.packages("psych")
```

## Generic library to use

```
library(ggplot2)
library(dplyr)
library(tidyverse)
library(readxl)
library(MASS)
library(car)
```

## Useful functions

```
trainindex = function(df, trainn = nrow(df)-nrow(df)/10){
  n = nrow(df)
  return(sample(1:n,trainn, replace=F))
}
```

## Data Reading

### Built in data

```
head(iris)
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
4	4.6	3.1	1.5	0.2	setosa
5	5.0	3.6	1.4	0.2	setosa
6	5.4	3.9	1.7	0.4	setosa

```
head(mtcars)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

### .txt file

```
cleaning = read.table("cleaning.txt", header=T)
head(cleaning)
```

	Case	Crews	Rooms	StdDev
1	1	16	51	12.000463
2	2	10	37	7.927123
3	3	12	37	7.289910
4	4	16	46	12.000463
5	5	16	45	12.000463
6	6	4	11	4.966555

### .csv file

```
Heart = read.csv("Heart.csv", header=T)
head(Heart)
```

	X	Age	Sex	ChestPain	RestBP	Chol	Fbs	RestECG	MaxHR	ExAng	Oldpeak	Slope
1	1	63	1	typical	145	233	1	2	150	0	2.3	3
2	2	67	1	asymptomatic	160	286	0	2	108	1	1.5	2
3	3	67	1	asymptomatic	120	229	0	2	129	1	2.6	2
4	4	37	1	nonanginal	130	250	0	0	187	0	3.5	3
5	5	41	0	nontypical	130	204	0	2	172	0	1.4	1
6	6	56	1	nontypical	120	236	0	0	178	0	0.8	1
Ca Thal AHD												
1	0		fixed	No								
2	3		normal	Yes								
3	2	reversable	Yes									
4	0		normal	No								
5	0		normal	No								

```
6 0      normal No
```

## .xls / .xlsx file

```
library(tidyverse)
library(readxl)
```

```
xlsxdata = read_excel("sample.xlsx")
xlsxdata
```

```
# A tibble: 9 x 4
  Age Gender Weight Height
<dbl> <chr>   <dbl>   <dbl>
1    14 Male     52    150
2    32 Male     82    167
3    54 Female    64    153
4    12 Male     45    147
5    71 Female    52    145
6    56 Female    65    161
7    23 Female    56    162
8    49 Male     92    186
9    30 Male     75    182
```

## .json file

## SQL

```
#library(dbplyr)
#library(RSQLite)
#dir.create("data", showWarnings = FALSE)
#download.file(url = "https://ndownloader.figshare.com/files/2292171", destfile = "data/portal_mammals.s
#mammals = DBI::dbConnect(RSQLite::SQLite(), "data/portal_mammals.sqlite")
```

```
#SQLdata = tbl(mammals, sql("SELECT year, species_id, plot_id FROM surveys"))
#head(SQLdata)
```

## Exploratory Analysis

Check List:

- General distribution of the object, widely spread? Skewed? Any outliers?
- Check for NA, NaN, or any other missing values equivalent.

## Summary

- Check for the basic summary: mean, median, min, max, frequency.
- For frequency, we prefer that all category data have sufficient number of data,
- Check for the correlation among variables.

```
summary(iris)
```

```

      Sepal.Length      Sepal.Width      Petal.Length      Petal.Width
Min.   :4.300      Min.   :2.000      Min.   :1.000      Min.   :0.100
1st Qu.:5.100      1st Qu.:2.800      1st Qu.:1.600      1st Qu.:0.300
Median :5.800      Median :3.000      Median :4.350      Median :1.300
Mean   :5.843      Mean   :3.057      Mean   :3.758      Mean   :1.199
3rd Qu.:6.400      3rd Qu.:3.300      3rd Qu.:5.100      3rd Qu.:1.800
Max.   :7.900      Max.   :4.400      Max.   :6.900      Max.   :2.500

      Species
setosa   :50
versicolor:50
virginica :50

```

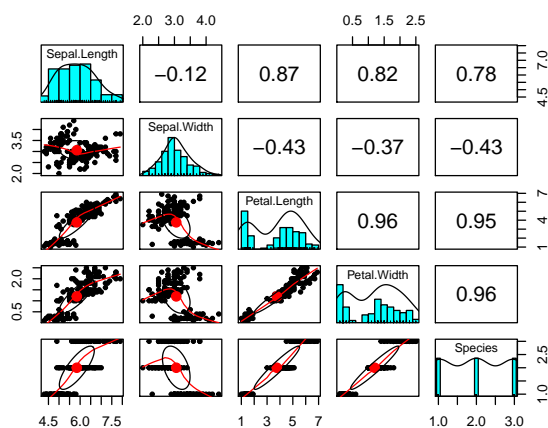
```
cor(iris[, -5])
```

```

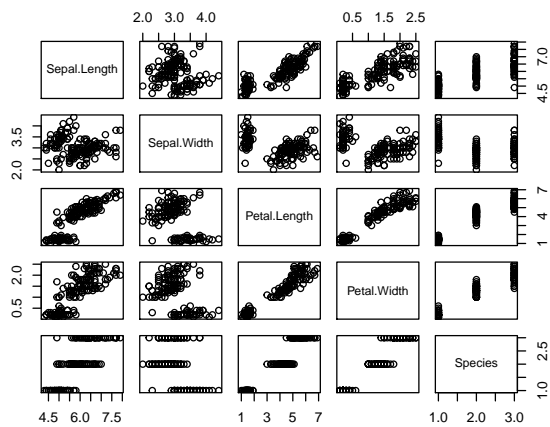
              Sepal.Length Sepal.Width Petal.Length Petal.Width
Sepal.Length    1.0000000  -0.1175698   0.8717538   0.8179411
Sepal.Width     -0.1175698   1.0000000  -0.4284401  -0.3661259
Petal.Length     0.8717538  -0.4284401   1.0000000   0.9628654
Petal.Width      0.8179411  -0.3661259   0.9628654   1.0000000

```

```
library(psych)
pairs.panels(iris)
```




```
pairs(iris)
```




## Histogram / Distributions


```
# easiest
hist(iris$Sepal.Length)
```




```
#
plot(density(iris$Sepal.Length), col="red")
```



```
# much easier for exploratory analysis
ggplot(data = iris) + geom_bar(mapping = aes(x = Sepal.Length))
```




```
ggplot(data=iris) + geom_histogram(mapping = aes(x=Sepal.Length))
```




```
ggplot(data=iris, mapping = aes(x=Sepal.Length, colour = Species)) + geom_freqpoly()
```






```
ggplot(Heart) + geom_density(mapping=aes(x=MaxHR, color=as.character(Sex)))
```



- Which values are most common among which category?
- Which values are rare, or odd? Could it be an outlier, or mis-interpreted?
- Any unusual patterns? Can you explain it?
- Why setosa tends to have smaller Sepal.Length than virginica?


## Boxplots

```
ggplot(data = iris) + geom_boxplot(mapping = aes(x=Species, y=Sepal.Length, fill=Species))
```




## Plots of two variables

```
ggplot(data=iris) + geom_point(aes(x=Sepal.Length, y=Sepal.Width, colour=Species))
```



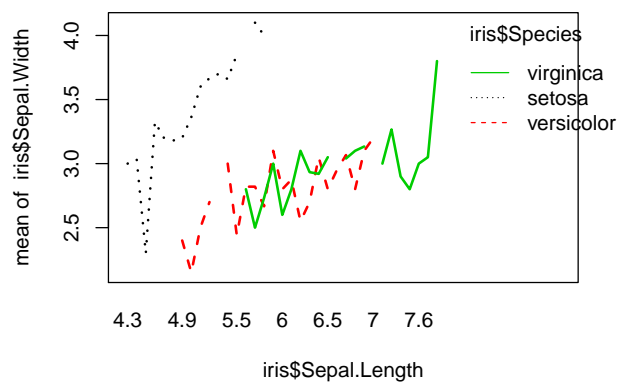
## Interaction Plots

```
ggplot(data=iris) + aes(x = Sepal.Length, y = Sepal.Width, colour = Species) +  
  geom_point(color = "grey") +  
  geom_smooth(method = "lm")
```



- Often used in time series

```
interaction.plot(x.factor = iris$Sepal.Length, trace.factor = iris$Species, response = iris$Sepal.Width
```



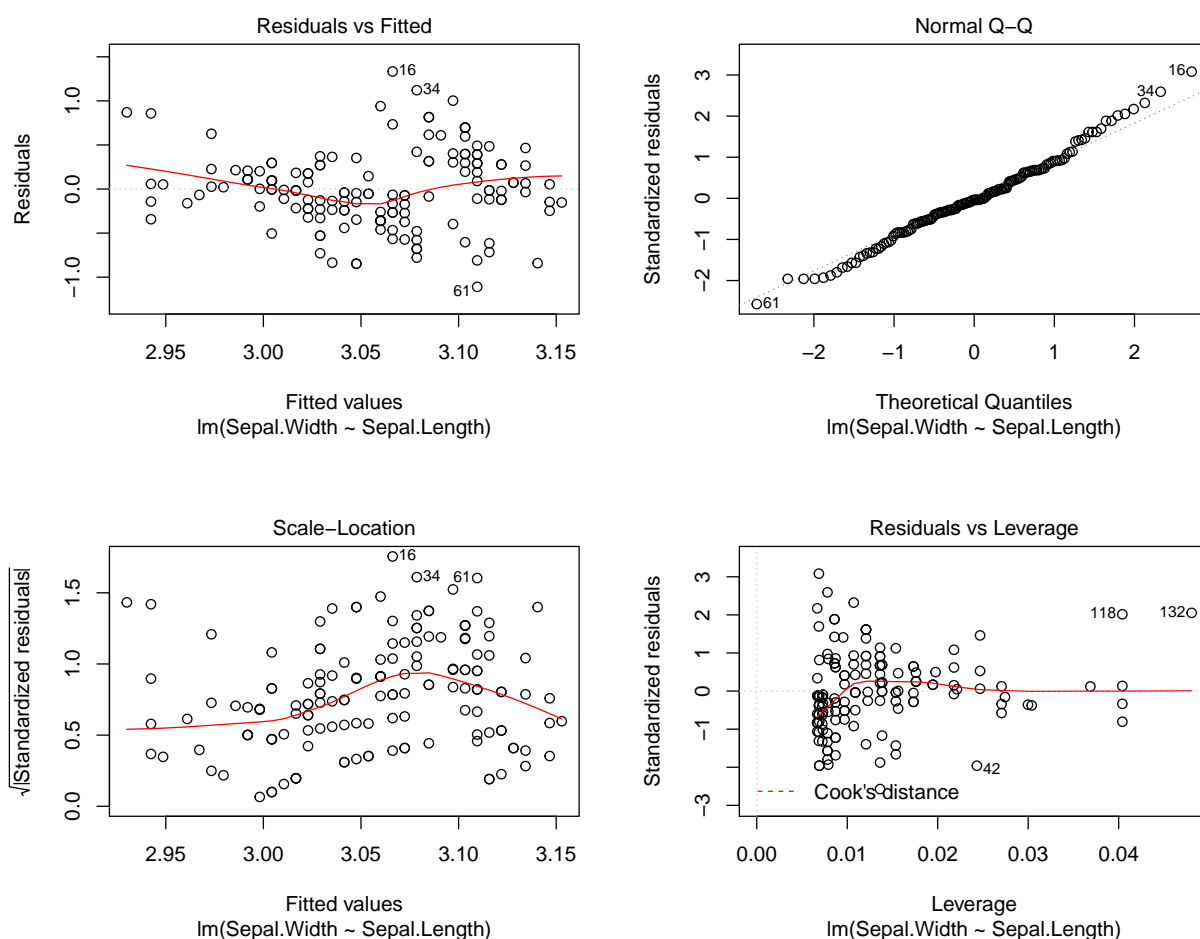
# Data Analysis, Prediction, and Classification

## Simple Linear Regression

### Assumption Check

- 1. Average is 0?
- 2. Standardized residual (more informative when leverage points exist because errors can show const var while residuals don't / how many estimated std deviations any point away from the fitted regression model / if outside -2 to 2, outlier), check for constant variance primarily here!
- 3. Normality (straight line) holds? For each  $x$ , see if corresponding  $y$  follow normal distributions where mean is fitted line.
- 4. Any outliers, leverage points (influential to fitted model / how predicted  $y$  change if removed / bigger than  $4/n$ ), outside Cook's distance ( $D_i = \frac{r_i^2}{2} * \frac{h_{ii}}{1-h_{ii}}$  /  $D_i > 4/(n-2)$ )?
- If any of the assumptions violated, any further inferences are invalidated.

```
model1 = lm(Sepal.Width~Sepal.Length, data=iris)
plot(model1)
```




```
head(lm.influence(model1)$hat)
```

```
1      2      3      4      5      6
0.012074844 0.015376584 0.019461346 0.021797361 0.013627836 0.008590398
```

## Statistical Inference


- Shape of the plots, linear? quadratic? exponential?

```
plot(Sepal.Width~Sepal.Length, data=iris);abline(model1)
```



*# Equivalently this, but much more complicated*

```
predicted1 = data.frame(Sepal.Length=iris$Sepal.Length, predicted = predict(model1, iris))
ggplot(data=iris) + geom_point(aes(x=Sepal.Length,y=Sepal.Width)) + geom_line(color='red',data = predicted1)
```



- Coefficients significant? P-value? Standard Error?
- Positive? Negative? Its strength?
- “One unit increase in X results in  $\beta_1$  much increase in Y”

$$\hat{\beta}_1 = r * \frac{S_Y}{S_X} = \frac{\sum^n (X - \bar{X})(Y - \bar{Y})}{\sum^n (Y - \bar{Y})^2}$$

$$\hat{\beta}_0 = \bar{Y} - \hat{\beta}_1 \bar{X}$$

$$e \sim N(0, \sigma^2)$$

$$S^2 = \frac{\sum^n (Y - \hat{Y})^2}{n-2}$$

$$se(\hat{\beta}_0)^2 = S^2 \left( \frac{1}{n} + \frac{(x^* - \bar{x})^2}{\sum^n (x - \bar{x})^2} \right)$$

$$se(\hat{\beta}_1)^2 = \frac{S^2}{\sum^n (X - \bar{X})^2}$$

- $SSE = \sum^n (y - \hat{y})^2$
- $SSR = \sum^n (\hat{y} - \bar{y})^2$
- $SST = \sum^n (y - \bar{y})^2$
- $R^2 = 1 - \frac{SSE}{SST}$
- $F = \frac{SST - SSE/1}{SST/n-2} = t^2 = \left(\frac{\hat{\beta}_1}{se(\hat{\beta}_1)}\right)^2$

```
summary(model1)
```

Call:

```
lm(formula = Sepal.Width ~ Sepal.Length, data = iris)
```

Residuals:

```
      Min       1Q   Median       3Q      Max
-1.1095 -0.2454 -0.0167  0.2763  1.3338
```

Coefficients:

```
              Estimate Std. Error t value Pr(>|t|)
(Intercept)   3.41895     0.25356   13.48  <2e-16 ***
Sepal.Length -0.06188     0.04297   -1.44    0.152
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Residual standard error: 0.4343 on 148 degrees of freedom

Multiple R-squared: 0.01382, Adjusted R-squared: 0.007159

F-statistic: 2.074 on 1 and 148 DF, p-value: 0.1519

```
var(iris$Sepal.Length)
```

```
[1] 0.6856935
```

- Confidence Interval:  $\hat{y} \pm t_{n-2} * S \sqrt{\frac{1}{n} + \frac{(x^* - \bar{x})^2}{\sum^n (x - \bar{x})^2}}$
- CI ex)  $\hat{y} \pm 2.06 * 0.4343 \sqrt{\frac{1}{150} + \frac{(5.0 - 5.843333)^2}{149 * 0.6856935}}$
- Prediction Interval:  $\hat{y} \pm t_{n-2} * S \sqrt{1 + \frac{1}{n} + \frac{(x^* - \bar{x})^2}{\sum^n (x - \bar{x})^2}}$
- PI ex)  $\hat{y} \pm 2.06 * 0.4343 \sqrt{1 + \frac{1}{150} + \frac{(5.0 - 5.843333)^2}{149 * 0.6856935}}$

## ANOVA SLR

Variation	df	SS	MS	F
Regression	1	SSR	SSR/1	SSR/(SSE/(n-2))
Residual/Error	n-2	SSE	SSE/n-2	
Total	n-1	SST		

```
anova(model1)
```

Analysis of Variance Table

Response: Sepal.Width

```
 Df Sum Sq Mean Sq F value Pr(>F)
```

---

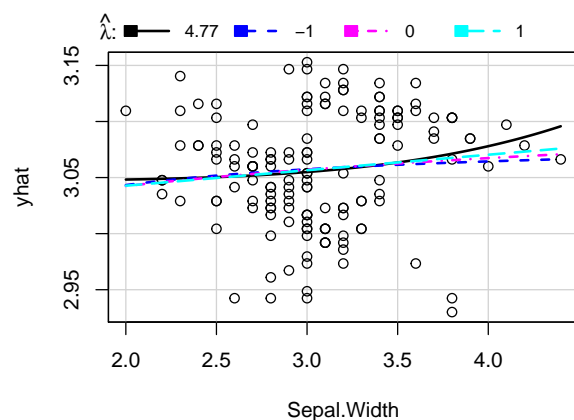
Sepal.Length	1	0.3913	0.39128	2.0744	0.1519
Residuals	148	27.9157	0.18862		

## Transformation

### Inverse response transformation

- $g^{-1}(y) = \beta_0 + \beta_1 x + \epsilon$ , e.g.  $g(y) = \exp(y)$ ,  $g^{-1}(y) = \log(y)$ ,  $g(y) = y^\lambda$ ,  $g^{-1}(y) = y^{1/\lambda}$ , make sure response only!
- Pick  $\lambda$  that has the lowest RSS/SSE.

```
library(alr3)
inverseResponsePlot(model1, key=TRUE)
```



	lambda	RSS
1	4.768992	0.3829388
2	-1.000000	0.3887806
3	0.000000	0.3873421
4	1.000000	0.3858685

### Box-cox transformation

- Try to make variables close to normally distributed. For SLR, maximize likelihood = minimize  $SSE(\lambda) = \sum (y^\lambda - \hat{\beta}_0 - \hat{\beta}_1 x)^2$ . Don't assume normality of  $x$ .
- Pick the  $\lambda = \text{Rounded Pwr}$

```
library(MASS)
model1bc = powerTransform(model1)
summary(model1bc)
```

```
bcPower Transformation to Normality
  Est Power Rounded Pwr Wald Lwr Bnd Wald Up Bnd
Y1    0.3429          1   -0.5634      1.2492
```

```
Likelihood ratio test that transformation parameter is equal to 0
(log transformation)
```

```
              LRT df    pval
LR test, lambda = (0) 0.5506328  1 0.45806
```

```
Likelihood ratio test that no transformation is needed
```

```
              LRT df    pval
LR test, lambda = (1) 2.01204  1 0.15606
```

### Log transformation



- Take logarithm on response or predictors or both.  $\log(y_2/y_1) = \beta_1 * \log(x_2/x_1)$ . “One percentage change in X results in  $\beta_1$  percentage change in Y”.
- Transform variables according to Rounded Pwr, e.g.  $Y1=\log(Y1)$ ,  $Y2=Y2^1$ .
- Rounded Pwr = 0 means log transform.
- e.g. if Rounded Pwr = 0.5 for Y3,  $Y3=\sqrt{Y3}$

```
model1pt = powerTransform(cbind(iris$Sepal.Length,iris$Sepal.Width)~1)
summary(model1pt)
```

bcPower Transformations to Multinormality

	Est	Power	Rounded Pwr	Wald Lwr Bnd	Wald Up Bnd
Y1	-0.1978		0	-1.2981	0.9026
Y2	0.3461		1	-0.5584	1.2506

Likelihood ratio test that transformation parameters are equal to 0  
(all log transformations)

	LRT	df	pval
LR test, lambda = (0 0)	0.6864385	2	0.70948

Likelihood ratio test that no transformations are needed

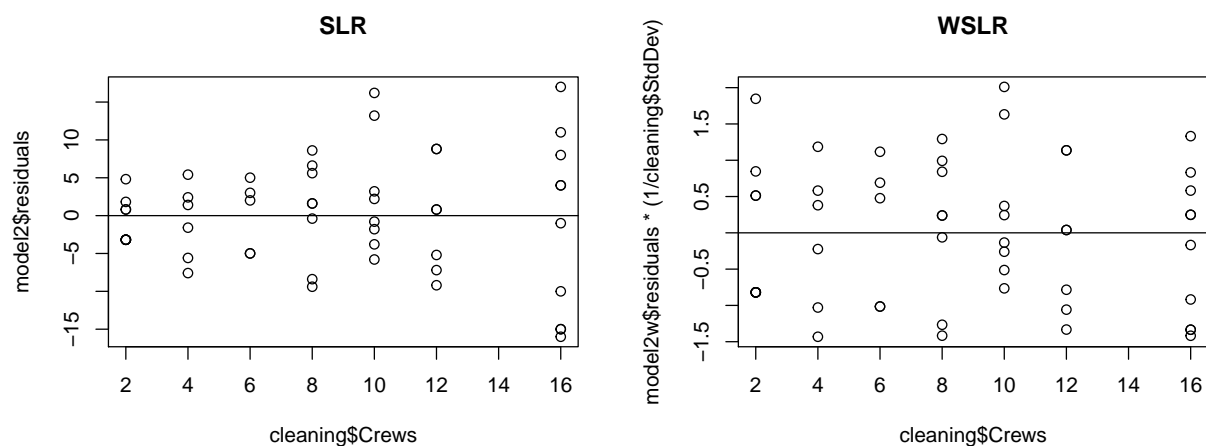
	LRT	df	pval
LR test, lambda = (1 1)	6.588234	2	0.037101

## Weighted Least Square

- When constant variance is violated, assigning a reasonable weight to each variance could fix the problem.
- Assign inversely proportional weights to the corresponding variances.
- $SSE = \sum^n (Y - (\hat{\beta}_0 + \hat{\beta}_1 X))^2$
- $WSSE = \sum^n w_i (Y - (\hat{\beta}_0 + \hat{\beta}_1 X))^2$  with  $\epsilon \sim N(0, \sigma^2/w_i)$
- Then,  $var(\sqrt{w_i}\epsilon_i) = \sigma^2$
- The weights are assumed to be known, so the estimated weights are used. Thus, this method works when the weights can be estimated precisely relative to one another.
- Sensitive to outliers, and possibly increase the influence of them.

```
model2 = lm(Rooms~Crews, data=cleaning)
model2w = lm(Rooms~Crews, weights=1/StdDev^2, data=cleaning)

plot(cleaning$Crews, model2$residuals, main="SLR");abline(h=0);plot(cleaning$Crews, model2w$residuals*(
```




## Multiple Linear Regression



### Assumptions

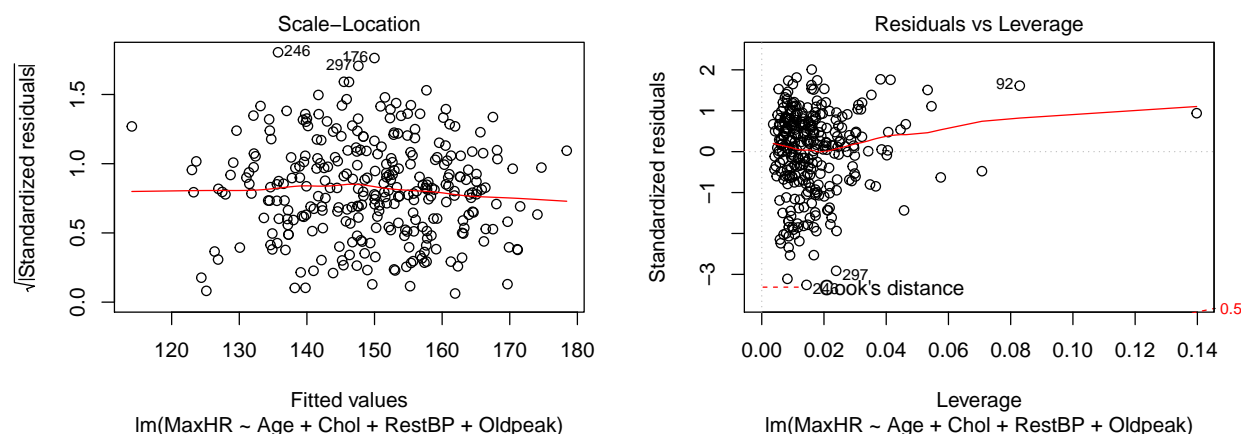
- The response variable and predictors are linearly related.
- Error terms are normally distributed.
- Error terms have a constant variance.
- Check the outliers, leverage points, and influential points.
- Check if the predictors are highly correlated (multicollinearity).
- Check the diagonal elements of the hat matrix  $\mathbf{H} = \mathbf{X}(\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T$ . If  $h_{ii} > 2 * \bar{h} = 2 * \frac{p+1}{n}$ , the point is considered to be a leverage point for MLR.
- Standardized residual  $r_i = \frac{e_i}{S\sqrt{1-h_{ii}}}$  where  $S^2 = \frac{SSE}{n-(p+1)}$ . If  $r_i$  is outside  $(-2, 2)$ , it's considered to be an outlier.
- Likewise, check for Cook's distance  $D_i$ . If greater than  $4/(n-2)$ , the point is an influential point for MLR.
- The diagnosis plots show if the entire model is valid.

```
m = lm(MaxHR~Age+Chol+RestBP+Oldpeak, data=Heart)
pairs(MaxHR~Age+Chol+RestBP+Oldpeak, data=Heart)
```



```
plot(m)
```





```
stdres1 = rstandard(m)
lev1 = hatvalues(m)
cookd1 = cooks.distance(m)
```

```
# outliers
which(abs(stdres1)>2)
```

```
24 38 80 88 115 172 176 224 229 237 245 246 253 297
24 38 80 88 115 172 176 224 229 237 245 246 253 297
```

```
# leverage ( hii > 2*(p+1)/n)
which(lev1 > 2*(4+1)/nrow(Heart))
```

```
4 49 69 92 114 122 124 127 153 162 182 184 189 192 202 212 232 274
4 49 69 92 114 122 124 127 153 162 182 184 189 192 202 212 232 274
286
286
```

```
# influential points (Di > 4/(n-2))
which(cookd1 > 4/(nrow(Heart)-2))
```

```
4 30 92 122 152 153 162 176 189 212 232 245 246 266 297
4 30 92 122 152 153 162 176 189 212 232 245 246 266 297
```

- When strong correlations exist among the predictor variables, the following issues may arise:
  1. F-test results will be highly significant, when very few predictors are significant.
  2. Some of the coefficients in the model show the opposite sign than expected.
- Variance Inflation Factor (VIF) :  $\frac{1}{1-R_j^2}$ , where  $R_j^2$  denote the value of  $R^2$  obtained from the regression of  $x_j$  on the other  $x$ 's. Note that  $var(\hat{\beta}_j) = \frac{1}{1-R_j^2} * \frac{\sigma^2}{(n-1)S_{x_j}^2}$ . If  $VIF_j > 5$ , then  $\hat{\beta}_j$  is poorly estimated due to multicollinearity.
- There are several ways to handle multicollinearity. One way is to delete the redundant predictors (highly correlated predictors).
- Another way to handle multicollinearity is to make the dataset uncorrelated i.e. linearly independent. The method is called Principle Component Analysis.

```
library(car)
vif(m)
```

```
Age      Chol    RestBP  Oldpeak
1.153910 1.051659 1.116389 1.064148
```

## Statistical Inference

- $\mathbf{Y} = \mathbf{X}\beta + \epsilon$
- $SSE = (\mathbf{Y} - \mathbf{X}\hat{\beta})^T(\mathbf{Y} - \mathbf{X}\hat{\beta}) = \|\mathbf{Y} - \mathbf{X}\hat{\beta}\|^2$
- $\hat{\beta} = (\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T\mathbf{Y}$
- $\mathbf{X} = (1, X_1, X_2, \dots, X_p)^T$
- $S^2 = \frac{SSE}{n-p-1}$
- $T_i = \frac{\hat{\beta}_i}{se(\hat{\beta}_i)} \sim t_{n-p-1}$  for  $H_0 : \beta_i = 0$
- $R^2 = 1 - \frac{SSE}{SST}$ , but always increase as  $p$  increase.
- $R_{adj}^2 = 1 - \frac{SSE/(n-p-1)}{SST/(n-1)}$
- $F = \frac{(SST-SSE)/p}{SSE/(n-p-1)} \sim F_{p,(n-p-1)}$  for  $H_0 : \beta_1 = \dots = \beta_p = 0$

```
summary(m)
```

Call:

```
lm(formula = MaxHR ~ Age + Chol + RestBP + Oldpeak, data = Heart)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-64.737	-11.889	3.353	13.875	39.861

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	181.31908	10.60124	17.104	< 2e-16 ***
Age	-0.96548	0.13690	-7.052	1.23e-11 ***
Chol	0.03343	0.02281	1.465	0.1439
RestBP	0.14069	0.06915	2.034	0.0428 *
Oldpeak	-5.70035	1.02343	-5.570	5.70e-08 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 20.02 on 298 degrees of freedom

Multiple R-squared: 0.2443, Adjusted R-squared: 0.2342

F-statistic: 24.09 on 4 and 298 DF, p-value: < 2.2e-16

# or

```
Y = Heart$MaxHR
```

```
X = cbind(1, Heart$Age, Heart$Chol, Heart$RestBP, Heart$Oldpeak)
```

```
beta = solve(t(X)%*%X)%*%t(X)%*%Y;beta
```

```



      [,1]
[1,] 181.31907819
[2,] -0.96547761
[3,]  0.03343059
[4,]  0.14068501
[5,] -5.70035136

```

## Model Diagnosis : Added Variable Plot

- Added variable plot enable us to visually assess the effect of each predictors, having adjusted for the effects of other predictors.
- In stead of the model  $\mathbf{Y} = \mathbf{X}\beta + \epsilon$ , consider  $\mathbf{Y} = \mathbf{X}\beta + \mathbf{Z}\alpha + \epsilon$ .
- $\mathbf{Z} = \mathbf{X}\delta + \epsilon$  and  $n \times 1$  vector. If the model with  $\mathbf{Z}$  fits better to the data, then the added variable plot should produce points randomly scattered around a line through thte origin with slope  $\hat{\alpha}$

```
library(car)
m = lm(iris$Sepal.Width~iris$Sepal.Length+iris$Petal.Length+iris$Petal.Width)
par(mfrow=c(2,2))
avPlot(m,variable=iris$Sepal.Length, ask=F)
avPlot(m,variable=iris$Petal.Length, ask=F)
avPlot(m,variable=iris$Petal.Width, ask=F)
par(mfrow=c(1,1))
```



## ANOVA MLR

Variation	df	SS	MS	F
Regression	p	SSR	SSR/p	(SST-SSE)/p/(SSE/(n-2))
Residual/Error	n-p-1	SSE	SSE/n-p-1	
Total	n-1	SST		

```
anova(model2)
```

## Analysis of Variance Table

Response: Rooms

```
      Df Sum Sq Mean Sq F value    Pr(>F)
Crews    1 16429.7  16429.7   305.27 < 2.2e-16 ***
Residuals 51  2744.8     53.8
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

## Nested Model and Partial F-test

$$H_o : \beta_1 = \dots = \beta_k = 0 \quad (k < p \text{ i.e. reduced model})$$

vs.

$$H_a : \beta_1 = \dots = \beta_p = 0 \quad (\text{i.e. full model})$$

- $F = \frac{SSE_{reduced} - SSE_{full} / (df_{reduced} - df_{full})}{SSE_{full} / df_{full}} = \frac{SSE_{reduced} - SSE_{full} / k}{SSE_{full} / df_{full}}$
- Although here the term “reduced” and “full” are used, make sure that “full” could already be a reduced model i.e.  $n - p - 1 \geq df_{full}$ .

```
model2reduced = lm(MaxHR~Age+RestBP+Oldpeak, data=Heart)
anova(model2,model2reduced)
```

## Analysis of Variance Table

Response: Rooms

```
      Df Sum Sq Mean Sq F value    Pr(>F)
Crews    1 16429.7  16429.7   305.27 < 2.2e-16 ***
Residuals 51  2744.8     53.8
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```


- Here, F-stats is not significant i.e. failed to reject null, reduced model is better.

## ANCOVA (Analysis of Covariance)

- Suppose we have a categorical variable with  $K$  levels. ANCOVA allows a categorical variable to be included in a linear model. Technically speaking, ANOVA is a sub-technique of ANCOVA where we have a hidden categorical variable with only 1 level.
- There could be multiple categorical variable. If that's the case, the number of levels would be  $K_1 \times K_2$ , and might require more sample size.
- Use variable selections (e.g. partial F-test) to reduce the number of predictors.
- $Y = \beta_0 + \beta_1 x + \beta_2 d + \beta_3 (d \times x) + \epsilon$  if  $d \in \{0, 1\}$
- $SSB = \sum_{i=1}^k \sum_{j=1}^{n_k} (\bar{Y}_i - \bar{\bar{Y}})^2$ , similar to SSE
- $SSW = \sum_{i=1}^k \sum_{j=1}^{n_k} (Y_{ij} - \bar{Y}_i)^2$ , similar to SSR
- $SST = \sum_{i=1}^k \sum_{j=1}^{n_k} (Y_{ij} - \bar{\bar{Y}})^2$
- $H_o : \mu_1 = \mu_2 = \dots = \mu_K$  vs.  $H_a$  : at least one of the group means is different

Variation	df	SS	MS	F
Between	k-1	SSB	SSB/(k-1)	SSB/(k-1)/(SSW/(n-k))
Within	n-k	SSW	SSW/n-k	
Total	n-1	SST		

```
ggplot(data = Heart) + geom_boxplot(mapping = aes(x=factor(Sex), y=MaxHR, fill=as.character(Sex)))
```



```
model3 = lm(MaxHR~Age*Sex, data=Heart)
# Same as lm(MaxHR~Age+Sex+Age:Sex)
summary(model3)
```

Call:

```
lm(formula = MaxHR ~ Age * Sex, data = Heart)
```

Residuals:

Min	1Q	Median	3Q	Max
-63.319	-13.192	3.475	15.630	46.337

Coefficients:

Estimate	Std. Error	t value	Pr(> t )
----------	------------	---------	----------




```
(Intercept) 198.9724    12.9278    15.391 < 2e-16 ***
Age          -0.8569     0.2288    -3.745 0.000217 ***
Sex           9.2725     15.7855     0.587 0.557373
Age:Sex       -0.2465     0.2827    -0.872 0.383945
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 21.01 on 299 degrees of freedom
Multiple R-squared:  0.1649,    Adjusted R-squared:  0.1565
F-statistic: 19.67 on 3 and 299 DF,  p-value: 1.147e-11
```

```
ggplot(data=Heart) + aes(x=Age, y=MaxHR, colour=as.character(Sex)) + geom_point() + geom_smooth(method="lm")
```



## Polynomial Regression

- $Y = \beta_0 + \beta_1 x + \beta_2 x^2 + \dots + \beta_h x^h + \epsilon$
- Pick the degree that gives the lowest MSE and highest R square adjusted.
- Be careful of over-fitting.

```
Y = iris$Sepal.Width
X = iris$Sepal.Length
polym1 = lm(Y~X)
summary(polym1)
```

Call:

```
lm(formula = Y ~ X)
```

Residuals:

Min	1Q	Median	3Q	Max
-1.1095	-0.2454	-0.0167	0.2763	1.3338

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	3.41895	0.25356	13.48	<2e-16 ***
X	-0.06188	0.04297	-1.44	0.152

---


Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4343 on 148 degrees of freedom

Multiple R-squared: 0.01382, Adjusted R-squared: 0.007159

F-statistic: 2.074 on 1 and 148 DF, p-value: 0.1519

```
qplot(X,polym1$fitted.values, geom=c("point","smooth"))
```



```
polym2 = lm(Y~X+I(X^2))
summary(polym2)
```

Call:

```
lm(formula = Y ~ X + I(X^2))
```

Residuals:

Min	1Q	Median	3Q	Max
-----	----	--------	----	-----

```
-1.13070 -0.26310 -0.02446  0.25728  1.38725
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	6.41584	1.58499	4.048	8.33e-05 ***
X	-1.08556	0.53625	-2.024	0.0447 *
I(X^2)	0.08571	0.04476	1.915	0.0574 .

---


Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4304 on 147 degrees of freedom

Multiple R-squared: 0.03783, Adjusted R-squared: 0.02474

F-statistic: 2.89 on 2 and 147 DF, p-value: 0.05877

```
qplot(X, polym2$fitted.values, geom=c("point", "smooth"))
```



```
polym3 = lm(Y~X+I(X^2)+I(X^3))
summary(polym3)
```

Call:

```
lm(formula = Y ~ X + I(X^2) + I(X^3))
```

Residuals:

Min	1Q	Median	3Q	Max
-1.17219	-0.23769	-0.00581	0.27359	1.34285

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-13.25795	9.98420	-1.328	0.1863
X	8.93524	5.05021	1.769	0.0789 .
I(X^2)	-1.58995	0.84097	-1.891	0.0607 .
I(X^3)	0.09202	0.04612	1.995	0.0479 *

---


Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4261 on 146 degrees of freedom

Multiple R-squared: 0.06337, Adjusted R-squared: 0.04412

F-statistic: 3.292 on 3 and 146 DF, p-value: 0.02239

```
qplot(X, polym3$fitted.values, geom=c("point", "smooth"))
```



```
polym4 = lm(Y~X+I(X^2)+I(X^3)+I(X^4))
summary(polym4)
```

Call:

```
lm(formula = Y ~ X + I(X^2) + I(X^3) + I(X^4))
```

Residuals:

	Min	1Q	Median	3Q	Max
	-1.24217	-0.23267	0.00282	0.25123	1.38777

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-142.56391	63.82905	-2.234	0.0270 *
X	97.48827	43.47380	2.242	0.0264 *
I(X^2)	-24.06445	10.99196	-2.189	0.0302 *
I(X^3)	2.59737	1.22267	2.124	0.0353 *
I(X^4)	-0.10351	0.05048	-2.051	0.0421 *

---


Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4215 on 145 degrees of freedom

Multiple R-squared: 0.08976, Adjusted R-squared: 0.06465

F-statistic: 3.575 on 4 and 145 DF, p-value: 0.008209

```
qplot(X, polym4$fitted.values, geom=c("point", "smooth"))
```



```
polym5 = lm(Y~X+I(X^2)+I(X^3)+I(X^4)+I(X^5))
summary(polym5)
```

Call:

```
lm(formula = Y ~ X + I(X^2) + I(X^3) + I(X^4) + I(X^5))
```

Residuals:

	Min	1Q	Median	3Q	Max
	-1.25194	-0.22446	-0.00225	0.25581	1.40416

Coefficients:


	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	-369.84591	416.99626	-0.887	0.377
X	292.79809	356.76489	0.821	0.413
I(X^2)	-90.55757	121.05337	-0.748	0.456
I(X^3)	13.80872	20.36289	0.678	0.499
I(X^4)	-1.03981	1.69825	-0.612	0.541
I(X^5)	0.03099	0.05618	0.552	0.582

Residual standard error: 0.4226 on 144 degrees of freedom

Multiple R-squared: 0.09168, Adjusted R-squared: 0.06014

F-statistic: 2.907 on 5 and 144 DF, p-value: 0.01568


```
qplot(X,polym5$fittd.values, geom=c("point","smooth"))
```



```
#MSEs
polyMSE = c(sqrt(sum(polym1$residuals)^2)/148,sqrt(sum(polym2$residuals)^2)/147,sqrt(sum(polym3$residuals)^2)/146,sqrt(sum(polym4$residuals)^2)/145,sqrt(sum(polym5$residuals)^2)/144)
polyMSE
```

```
[1] 8.532964e-18 7.127707e-18 3.707080e-17 3.825963e-17 2.650272e-18
```

```
plot(1:5, polyMSE, type="b", main="degree vs residual", ylab="MSE")
```



## Model Selection Criteria

- When multicollinearity among the predictor variables are observed, model selection is one way to resolve it.
- Make sure to check all the assumptions are met before starting model selection.
- Goodness of fit criteria:
  1. Adjusted R-square:  $R_{adj}^2 = 1 - \frac{SSE/(n-p-1)}{SST/(n-1)} = 1 - \frac{(1-R^2)(n-1)}{(n-p-1)}$ . We don't use R-square because it automatically increase as the number of predictors increase.
  2. Akaike Information Criterion (AIC): Smaller the better. Reward for a good fit + penalty for complexity.  $AIC = n * \log(\frac{SSE}{n}) + 2p$
  3. AIC corrected: Greater penalty. Smaller the better.  $AIC_C = AIC + \frac{2p(p+2)(p+3)}{n-p-1}$
  4. Bayes Information Criteria (BIC): Greater penalty than AIC when  $\log(n) > 2$ , thus favors simpler model than  $AIC$ . As the sample size  $n$  increase, the probability that BIC choose the correct model becomes 1. For smaller  $n$ , BIC choose too simple model hence biased.  $BIC = n * \log(\frac{SSE}{n}) + \log(n) * p$
  5. Mean Square Error: When test data is given, MSE is the most reliable measurement for choosing the best model. Combine with CV.

## Subset Selection: Best Subset Model

- For  $p$  predictors, we have  $\sum_{k=1}^p \binom{p}{k}$  possible subset models.
  1. Start with  $k = p$  i.e. full model, fit the model.
  2.  $k = p - 1$ . Fit all  $\binom{p}{p-1}$  models, keep the winner among  $\binom{p}{p-1}$  with highest  $R^2$ . Here it's  $R^2$  because we are comparing the models with same number of predictors.
  3.  $k = p - 2$ , keep the winner.
  4. Repeat until  $k = 1$ .
  5. Choose the best model among all winners. Use the criteria other than  $R^2$ .

## R: Best Subset Model

```
library(leaps)
X = cbind(iris$Sepal.Length, iris$Petal.Length, iris$Petal.Width)
b = regsubsets(X, iris$Sepal.Width)
bs = summary(b);bs
```

```
Subset selection object
3 Variables (and intercept)
  Forced in Forced out
a      FALSE      FALSE
b      FALSE      FALSE
c      FALSE      FALSE
1 subsets of each size up to 3
Selection Algorithm: exhaustive
      a      b      c
1 ( 1 ) " " "*" " "
2 ( 1 ) "*" "*" " "
3 ( 1 ) "*" "*" "*"
```

```

om1 = lm(Sepal.Width~Petal.Length,data=iris)
om2 = lm(Sepal.Width~Sepal.Length+Petal.Length,data=iris)
om3 = lm(Sepal.Width~Sepal.Length+Petal.Length+Petal.Width,data=iris)

n = nrow(iris)

p=1
AIC1 = extractAIC(om1,k=2)[2]
AICc1 = extractAIC(om1,k=2)[2] + 2*(p+2)*(p+3)/(n-p-1)
BIC1 = extractAIC(om1, k=log(n))[2]

p=2
AIC2 = extractAIC(om1,k=2)[2]
AICc2 = extractAIC(om1,k=2)[2] + 2*(p+2)*(p+3)/(n-p-1)
BIC2 = extractAIC(om1, k=log(n))[2]

p=3
AIC3 = extractAIC(om1,k=2)[2]
AICc3 = extractAIC(om1,k=2)[2] + 2*(p+2)*(p+3)/(n-p-1)
BIC3 = extractAIC(om1, k=log(n))[2]

AIC = c(AIC1,AIC2,AIC3)
AICc = c(AICc1, AICc2, AICc3)
BIC = c(BIC1, BIC2, BIC3)

data.frame(Radj2 = bs$adjr2,AIC,AICc,BIC)

```

	Radj2	AIC	AICc	BIC
1	0.1780444	-276.5497	-276.3876	-270.5285
2	0.4490193	-276.5497	-276.2776	-270.5285
3	0.5142264	-276.5497	-276.1388	-270.5285

---



**R: Best Subset Model (Alternative)**

- Black represents that the variables are selected. Pick the variables that highlighted on the very top.

```
library(ISLR)
Hitters=na.omit(Hitters)
regfit.full=regsubsets(Salary~.,Hitters)
summary(regfit.full)
```

Subset selection object

Call: regsubsets.formula(Salary ~ ., Hitters)

19 Variables (and intercept)

Forced in Forced out

AtBat	FALSE	FALSE
Hits	FALSE	FALSE
HmRun	FALSE	FALSE
Runs	FALSE	FALSE
RBI	FALSE	FALSE
Walks	FALSE	FALSE
Years	FALSE	FALSE
CAtBat	FALSE	FALSE
CHits	FALSE	FALSE
CHmRun	FALSE	FALSE
CRuns	FALSE	FALSE
CRBI	FALSE	FALSE
CWalks	FALSE	FALSE
LeagueN	FALSE	FALSE
DivisionW	FALSE	FALSE
PutOuts	FALSE	FALSE
Assists	FALSE	FALSE
Errors	FALSE	FALSE
NewLeagueN	FALSE	FALSE

1 subsets of each size up to 8


Selection Algorithm: exhaustive

		AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAtBat	CHits	CHmRun	CRuns
1	( 1 )	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "	" "
2	( 1 )	" "	"*	" "	" "	" "	" "	" "	" "	" "	" "	" "
3	( 1 )	" "	"*	" "	" "	" "	" "	" "	" "	" "	" "	" "
4	( 1 )	" "	"*	" "	" "	" "	" "	" "	" "	" "	" "	" "
5	( 1 )	"*	"*	" "	" "	" "	" "	" "	" "	" "	" "	" "
6	( 1 )	"*	"*	" "	" "	" "	"*	" "	" "	" "	" "	" "
7	( 1 )	" "	"*	" "	" "	" "	"*	" "	"*	"*	"*	" "
8	( 1 )	"*	"*	" "	" "	" "	"*	" "	" "	" "	"*	"*


  

		CRBI	CWalks	LeagueN	DivisionW	PutOuts	Assists	Errors	NewLeagueN
1	( 1 )	"*	" "	" "	" "	" "	" "	" "	" "
2	( 1 )	"*	" "	" "	" "	" "	" "	" "	" "
3	( 1 )	"*	" "	" "	" "	"*	" "	" "	" "
4	( 1 )	"*	" "	" "	"*	"*	" "	" "	" "
5	( 1 )	"*	" "	" "	"*	"*	" "	" "	" "
6	( 1 )	"*	" "	" "	"*	"*	" "	" "	" "
7	( 1 )	" "	" "	" "	"*	"*	" "	" "	" "
8	( 1 )	" "	"*	" "	"*	"*	" "	" "	" "

```
par(mfrow=c(1,1))
plot(regfit.full,scale="adjr2")
```



```
plot(regfit.full,scale="bic")
```



## Subset Selection: Stepwise Regression

### Forward Stepwise

1. Start with null model (intercept only)
  2. Fit model with  $k = 1$ . Choose the best among  $p$  models based on  $R^2$ .
  3. Add another variable to the previous model, keep the best model.
  4. Repeat until  $k = p$ .
  5. Choose the best among the  $p$  candidate models using AIC, BIC etc...
- 

### R: Forward AIC

```
m = lm(Sepal.Width~1, data=iris)

forwardAIC = step(m, scope = list(lower=~1,
upper=~Sepal.Length+Petal.Length+Petal.Width,data),
direction = "forward", data=iris)
```

Start: AIC=-248.13

Sepal.Width ~ 1

	Df	Sum of Sq	RSS	AIC
+ Petal.Length	1	5.1960	23.111	-276.55
+ Petal.Width	1	3.7945	24.512	-267.72
+ Sepal.Length	1	0.3913	27.916	-248.22
<none>			28.307	-248.13

Step: AIC=-276.55

Sepal.Width ~ Petal.Length

	Df	Sum of Sq	RSS	AIC
+ Sepal.Length	1	7.7237	15.387	-335.56
+ Petal.Width	1	0.8363	22.275	-280.08
<none>			23.111	-276.55

Step: AIC=-335.56

Sepal.Width ~ Petal.Length + Sepal.Length

	Df	Sum of Sq	RSS	AIC
+ Petal.Width	1	1.9133	13.474	-353.48
<none>			15.387	-335.56

Step: AIC=-353.48

Sepal.Width ~ Petal.Length + Sepal.Length + Petal.Width

**R: Forward BIC**

```
forwardBIC = step(m, scope = list(lower=~1,
upper=~Sepal.Length+Petal.Length+Petal.Width,data),
direction = "forward", data=iris, k=log(n))
```

Start: AIC=-245.12

Sepal.Width ~ 1

	Df	Sum of Sq	RSS	AIC
+ Petal.Length	1	5.1960	23.111	-270.53
+ Petal.Width	1	3.7945	24.512	-261.70
<none>			28.307	-245.12
+ Sepal.Length	1	0.3913	27.916	-242.20

Step: AIC=-270.53

Sepal.Width ~ Petal.Length

	Df	Sum of Sq	RSS	AIC
+ Sepal.Length	1	7.7237	15.387	-326.53
+ Petal.Width	1	0.8363	22.275	-271.05
<none>			23.111	-270.53

Step: AIC=-326.53

Sepal.Width ~ Petal.Length + Sepal.Length

	Df	Sum of Sq	RSS	AIC
+ Petal.Width	1	1.9133	13.474	-341.44
<none>			15.387	-326.53

Step: AIC=-341.44

Sepal.Width ~ Petal.Length + Sepal.Length + Petal.Width

## Backward Stepwise

1. Start with full model.
  2. Fit all models with  $k = p - 1$ , pick the best model.
  3. Reduce another variable from 2., pick the best model.
  4. Repeat until  $k = 1$ .
  5. Choose the best among the  $p$  candidate models using AIC, BIC etc...
- 

## R: Backward AIC

```
m = lm(mpg~disp+hp+drat+wt+qsec,data=mtcars)
backAIC = step(m,direction = "backward", data=mtcars)
```

Start: AIC=65.47

mpg ~ disp + hp + drat + wt + qsec

	Df	Sum of Sq	RSS	AIC
- disp	1	3.974	174.10	64.205
<none>			170.13	65.466
- hp	1	11.886	182.01	65.627
- qsec	1	12.708	182.84	65.772
- drat	1	15.506	185.63	66.258
- wt	1	81.394	251.52	75.978

Step: AIC=64.21

mpg ~ hp + drat + wt + qsec

	Df	Sum of Sq	RSS	AIC
- hp	1	9.418	183.52	63.891
- qsec	1	9.578	183.68	63.919
<none>			174.10	64.205
- drat	1	11.956	186.06	64.331
- wt	1	113.882	287.99	78.310

Step: AIC=63.89

mpg ~ drat + wt + qsec

	Df	Sum of Sq	RSS	AIC
<none>			183.52	63.891
- drat	1	11.942	195.46	63.908
- qsec	1	85.720	269.24	74.156
- wt	1	275.686	459.21	91.241

## R: Backward BIC

```
backBIC = step(m, direction="backward", data=mtcars, k=log(n))
```

Start: AIC=83.53

mpg ~ disp + hp + drat + wt + qsec

	Df	Sum of Sq	RSS	AIC
- disp	1	3.974	174.10	79.258
- hp	1	11.886	182.01	80.681
- qsec	1	12.708	182.84	80.825
- drat	1	15.506	185.63	81.311
<none>			170.13	83.530
- wt	1	81.394	251.52	91.031

Step: AIC=79.26

mpg ~ hp + drat + wt + qsec

	Df	Sum of Sq	RSS	AIC
- hp	1	9.418	183.52	75.934
- qsec	1	9.578	183.68	75.962
- drat	1	11.956	186.06	76.373
<none>			174.10	79.258
- wt	1	113.882	287.99	90.352

Step: AIC=75.93

mpg ~ drat + wt + qsec

	Df	Sum of Sq	RSS	AIC
- drat	1	11.942	195.46	72.940
<none>			183.52	75.934
- qsec	1	85.720	269.24	83.188
- wt	1	275.686	459.21	100.272

Step: AIC=72.94

mpg ~ wt + qsec

	Df	Sum of Sq	RSS	AIC
<none>			195.46	72.940
- qsec	1	82.86	278.32	79.239
- wt	1	733.19	928.66	117.797

## R: Forward and Backward Selection (Alternative)


```
regfit.fwd=regsubsets(Salary~.,data=Hitters,nvmax=19,method="forward")
regfit.bwd=regsubsets(Salary~.,data=Hitters,nvmax=19,method="backward")

# Choosing best number of predictors w/ CV, based on MSE.

k=10
set.seed(1)
folds=sample(1:k,nrow(Hitters),replace=TRUE)
cv.errors=matrix(NA,k,19, dimnames=list(NULL, paste(1:19)))

predict.regsubsets=function(object,newdata,id,...){
  form=as.formula(object$call[[2]])
  mat=model.matrix(form,newdata)
  coefi=coef(object,id=id)
  xvars=names(coefi)
  mat[,xvars]%*%coefi
}

for(j in 1:k){
  best.fit=regsubsets(Salary~.,data=Hitters[folds!=j,],nvmax=19)
  for(i in 1:19){
    pred=predict(best.fit,Hitters[folds==j,],id=i)
    cv.errors[j,i]=mean((Hitters$Salary[folds==j]-pred)^2)
  }
}
mean.cv.errors=apply(cv.errors,2,mean)
plot(mean.cv.errors,type='b')
```



```
reg.best=regsubsets(Salary~.,data=Hitters, nvmax=19);coef(reg.best,11)
```

(Intercept)	AtBat	Hits	Walks	CAtBat
135.7512195	-2.1277482	6.9236994	5.6202755	-0.1389914
CRuns	CRBI	CWalks	LeagueN	DivisionW
1.4553310	0.7852528	-0.8228559	43.1116152	-111.1460252
PutOuts	Assists			
0.2894087	0.2688277			

## Shrinkage: Regularization

- Fit a model involving all  $p$  predictors, but the estimated coefficients are shrunk toward zero. This shrinkage has the effect of reducing variance and can also perform variable selection.
- For both shrinkage methods below, it is recommended to use CV to choose the best tuning parameter.

## Ridge Regression

- Instead of minimizing loss function  $SSE = (Y - X\hat{\beta})^T(Y - X\hat{\beta})$ , minimize the loss function with L2 penalty term:
- Minimize  $(Y - X\hat{\beta})^T(Y - X\hat{\beta}) + \lambda\hat{\beta}^T\hat{\beta} = SSE + \lambda\|\hat{\beta}\|_2^2$  where  $\lambda \geq 0$
- The coefficients generated by Ridge regression tends to be similar in value (absolute value-wise) as the tuning parameter  $\lambda$  increases. This is because the penalty term uses Euclidean distance, or L2 norm. When the tuning parameter is too large, all coefficients goes down to 0.
- Ridge regression coefficient estimate can change substantially when multiplying a given predictor by a constant, due to L2 norm part. Thus, it is best to apply Ridge regression after standardizing the predictors using the formula:  $\tilde{x}_{ij} = \frac{x_{ij}}{\sqrt{\frac{1}{n} \sum_{i=1}^n (x_{ij} - \bar{x}_j)^2}}$

## LASSO regression

- Minimize  $(Y - X\hat{\beta})^T(Y - X\hat{\beta}) + \lambda\mathbf{1}^T|\hat{\beta}| = SSE + \lambda\|\hat{\beta}\|_1$  where  $\lambda \geq 0$
- The coefficients generated by LASSO regression tends to be reduced down to 0 as the tuning parameter increases. This is because the penalty term uses Manhattan distance, or L1 norm. Thus, this shrinkage method can perform as variable selection.
- We say that the lasso yeilds sparse models i.e. the models that involve only a subset of variables, or contains many zeros.

## Elastic Net

- The combination of LASSO and Ridge.
- Minimize  $SSE + \lambda(\frac{1-\alpha}{2}\|\hat{\beta}\|_2^2 + \alpha\|\hat{\beta}\|_1)$  where  $\lambda \geq 0$ ,  $\alpha \in (0, 1)$



**R: Shrinkage**

- Ridge:  $\alpha = 0$
- LASSO:  $\alpha = 1$
- Elastic Net:  $\alpha \in (0, 1)$
- Make sure to use matrix / vector for the inputs. Sparse matrix from `library(Matrix)` is supported.

**R: Ridge Regression**

```
library(glmnet)
library(ISLR) # For Hitters dataset
x=model.matrix(Salary~.,Hitters)[,-1]
y=Hitters$Salary

set.seed(1)
train=trainindex(x)
temp = 1:nrow(x)
test=temp[-train]
y.test=y[test]
grid=10^seq(10,-2,length=100)
ridge.mod=glmnet(x[train,],y[train], alpha=0,lambda=grid, thresh=1e-12)

# lambda = 4
ridge.pred=predict(ridge.mod,s=4,newx=x[test,])
mean((ridge.pred-y.test)^2)

[1] 165161.7

# lambda = 0
ridge.pred=predict(ridge.mod,s=0,newx=x[test,])
mean((ridge.pred-y.test)^2)

[1] 163360

# When lambda = 0, same as lm
lm(y~x, subset=train)
```

Call:

```
lm(formula = y ~ x, subset = train)
```

Coefficients:

(Intercept)	xAtBat	xHits	xHmRun	xRuns
219.6948	-2.1459	7.4588	0.3843	-2.5243
xRBI	xWalks	xYears	xCAtBat	xCHits
0.6735	6.6994	-11.9238	-0.0865	0.1849
xCHmRun	xCRuns	xCRBI	xCWalks	xLeagueN
-0.2475	0.9189	0.9029	-0.9902	68.1354
xDivisionW	xPutOuts	xAssists	xErrors	xNewLeagueN
-145.4021	0.3320	0.3830	-5.3149	-55.1098

```
ridge.mod=glmnet(x[train,],y[train], alpha=0,lambda=0, thresh=1e-12)
predict(ridge.mod,s=0,exact=T,type="coefficients")
```

20 x 1 sparse Matrix of class "dgCMatrix"

```


1
(Intercept) 219.69034188
AtBat       -2.14582772
Hits        7.45829955
HmRun       0.38382694
Runs       -2.52386914
RBI         0.67358027
Walks       6.69923354
Years     -11.92178980
CAtBat     -0.08654625
CHits      0.18518313
CHmRun     -0.24717482
CRuns      0.91871110
CRBI       0.90278683
CWalks    -0.99010232
LeagueN   68.13853859
DivisionW -145.40170447
PutOuts    0.33202384
Assists    0.38306254
Errors    -5.31509262
NewLeagueN -55.11307800

```

```

# Cross Validation to choose lambda
cv.out=cv.glmnet(x[train,],y[train],alpha=0)
plot(cv.out)

```



```

bestlam=cv.out$lambda.min
bestlam

```

```
[1] 25.58362
```

```


ridge.mod=glmnet(x[train,],y[train], alpha=0,lambda=bestlam, thresh=1e-12)
ridge.pred=predict(ridge.mod, newx=x[test,])
mean((ridge.pred-y.test)^2)

```


```
[1] 149900.7
```

R: LASSO Regression

```
lasso.mod=glmnet(x[train,],y[train],alpha=1,lambda=grid)
plot(lasso.mod)
```



```
set.seed(1)
cv.out=cv.glmnet(x[train,],y[train],alpha=1)
plot(cv.out)
```



```
bestlam=cv.out$lambda.min
lasso.mod=glmnet(x[train,],y[train],alpha=1,lambda=bestlam)
lasso.pred=predict(lasso.mod,newx=x[test,])
mean((lasso.pred-y.test)^2)
```

[1] 163253

```
out=glmnet(x,y,alpha=1,lambda=grid)
lasso.coef=predict(out,type="coefficients",s=bestlam)[1:20,];lasso.coef[lasso.coef!=0]
```

(Intercept)	AtBat	Hits	HmRun	Runs
149.54920606	-1.88397299	6.55952778	0.88197054	-0.59100530
Walks	Years	CAAtBat	CHmRun	CRuns
5.43552977	-8.55446820	-0.04530057	0.34497437	1.00059582
CRBI	CWalks	LeagueN	DivisionW	PutOuts
0.51183549	-0.69430903	42.45335427	-117.37199676	0.28122999

---

Assists	Errors	NewLeagueN
0.26777558	-2.71817199	-6.60210398

## Dimension Reduction: Principle Component Analysis

- Project the  $p$  predictors into  $M$ -dimensional subspace, where  $M < p$ . This is achieved by computing  $M$  different linear combinations of the variables. Then,  $M$  projections are used as predictor to fit a linear regression model by least squares.
- Before applying PCA, normalize the dataset.
- The first principle component is the linear combination of the variables with the largest variance so that the dataset is easily distinguishable e.g. LDA or grouping would perform well.
- The second principle component is also the linear combination of the variables with the largest variance, subject to being uncorrelated with the first one i.e. perpendicular / independent of the first one, and third PC, forth PC, and so on.
- Drawback: the directions / new dimensions identified by PCA, i.e. the way linear combinations are created, is unsupervised way since the response  $Y$  is not used to help determine the PCA directions. Thus, there is no gurantee that the directions that best explain the predictors will also be the best directions to use for predicting the response.
- Regression on data with PCA applied is called principle component regression.
- Mathematically, we have the following optimization problem. Solving this using the method of Lagrange multiplier, it reveals that the PCA coefficients of the linear combinations are the eigenvectors of the variance-covariance matrix of a dataset with the eigenvalues equal to the Lagrange multipliers.
- A dataset with all PC's applied has a variance-covariance matrix whose non-diagonal entires are zero. That is, the all variables are uncorrelated.

$$\min \mathbf{W}^T \Sigma \mathbf{W} \text{ s.t. } \mathbf{W}^T \mathbf{W} = 1$$

### R: PCA

```
mtcars_temp = subset(mtcars, select=c(mpg, wt, disp, qsec))
mtcars_c = scale(mtcars_temp) # Normalize
out_pca = princomp(mtcars_c) # PCA

summary(out_pca)
```

Importance of components:

	Comp.1	Comp.2	Comp.3	Comp.4
Standard deviation	1.681763	0.9178832	0.3865919	0.23390037
Proportion of Variance	0.729891	0.2174218	0.0385686	0.01411855
Cumulative Proportion	0.729891	0.9473129	0.9858814	1.00000000

```
W = out_pca$loadings; W # Linear combinations
```

Loadings:

	Comp.1	Comp.2	Comp.3	Comp.4
mpg	0.555		0.762	0.327
wt	-0.538	-0.376	0.105	0.747
disp	-0.561		0.639	-0.523
qsec	0.297	-0.922		-0.247

	Comp.1	Comp.2	Comp.3	Comp.4
SS loadings	1.00	1.00	1.00	1.00
Proportion Var	0.25	0.25	0.25	0.25
Cumulative Var	0.25	0.50	0.75	1.00

```
# Dimension Reduction to PC1
```

```
X = as.matrix(mtcars_c)
```

```
XW = X%*%W ; XW[,1]
```

Mazda RX4	Mazda RX4 Wag	Datsun 710
0.5012988	0.4542238	1.4243651
Hornet 4 Drive	Hornet Sportabout	Valiant
0.2627126	-0.9732059	0.1029417
Duster 360	Merc 240D	Merc 230
-1.6460261	1.1495838	1.5324653
Merc 280	Merc 280C	Merc 450SE
0.1561807	0.1268289	-1.0870481
Merc 450SL	Merc 450SLC	Cadillac Fleetwood
-0.7841288	-0.9386639	-3.0797248
Lincoln Continental	Chrysler Imperial	Fiat 128
-3.1475747	-2.6838810	2.6502919
Honda Civic	Toyota Corolla	Toyota Corona
2.6431662	3.0948242	1.4026790
Dodge Challenger	AMC Javelin	Camaro Z28
-1.1468472	-0.9930228	-1.9126488
Pontiac Firebird	Fiat X1-9	Porsche 914-2
-1.3256578	2.2299109	1.4452233
Lotus Europa	Ford Pantera L	Ferrari Dino
2.3423493	-1.4698954	0.2075584
Maserati Bora	Volvo 142E	
-1.5204547	0.9821761	

```
# Covariance Matrix
```

```
var(X)
```

	mpg	wt	disp	qsec
mpg	1.0000000	-0.8676594	-0.8475514	0.4186840
wt	-0.8676594	1.0000000	0.8879799	-0.1747159
disp	-0.8475514	0.8879799	1.0000000	-0.4336979
qsec	0.4186840	-0.1747159	-0.4336979	1.0000000

```
round(var(XW),7)
```

	Comp.1	Comp.2	Comp.3	Comp.4
Comp.1	2.919564	0.0000000	0.0000000	0.0000000
Comp.2	0.000000	0.8696873	0.0000000	0.0000000
Comp.3	0.000000	0.0000000	0.1542744	0.0000000
Comp.4	0.000000	0.0000000	0.0000000	0.0564742


## R: Principle Component Regression

- Pick ncomp that minimize MSEP.

```
library(pls)
train.i = trainindex(mtcars_temp)

mtcars_train = mtcars_temp[train.i,]
mtcars_test = mtcars_temp[-train.i,2:4]
y_test = mtcars_temp[-train.i,1]

pcr.fit=pcr(mpg~., data=mtcars_train,scale=TRUE, validation="CV")
validationplot(pcr.fit,val.type="MSEP")
```



```
pcr.pred=predict(pcr.fit,mtcars_test,ncomp=1)
mean((pcr.pred-y_test)^2)
```

```
[1] 8.623331
```

```
pcr.fit=pcr(y~x,scale=TRUE,ncomp=1)
summary(pcr.fit)
```

```
Data:   X dimension: 263 19
       Y dimension: 263 1
Fit method: svdpc
Number of components considered: 1
TRAINING: % variance explained
  1 comps
X    38.31
y    40.63
```

## Spline

- The truth is that association between variables are hardly ever linear. However, polynomial regression often cannot capture all features and associations among variables. To achieve this, fit linear regression piece-wise, and fit different polynomials for each pieces, then smooth the entire curve. Thus, spline works on one continuous predictor.
- Keep bias-variance trade off in mind. Linear model is highly biased, but if we attempt to increase the flexibility of the model, the bias decreases but variability increases. Add more flexibility until the bias is low enough yet the variability is not too high.
- The junctions of pieces is called “knot”.
- Step functions are used to divide the predictor variables and add knots. `cut` command divides the predictor variables into K many equal length intervals and assign each observation into the appropriate interval. Use `break` option to make your own cut if desired.
- Spline is the combine step functions and polynomial regression. One draw back is that piecewise polynomials can create a discontinuity at knots. Thus, we smooth the curve.
- `bs(X, knot)`: Any degree splines e.g. Linear splines and cubic splines (piece-wise cubic polynomials). A cubic spline with K knots has K+4 df.
- `ns(X, df)`: Natural cubic splines. A natural cubic spline extrapolates linearly beyond the boundary knots. This adds 4 extra constraints, and allows us to put more internal knots for the same degrees of freedom as a regular cubic spline. Thus, natural splines with K knots has K df.
- `s(X, df)`: Smoothing splines, fit model via minimize  $SSE + \lambda \int g''(t)dt$ . Smoothing splines avoid the knot-selection issue, leaving a single  $\lambda$  to be chosen. Here, `smooth.spline()` function fit a smoothing spline, and we choose df instead of  $\lambda$ .
- `lo(X, span)`: Local regression spline, also called loess. We fit separate linear fits over the range of the predictor variable by weighted least squares. Highly flexible. Use `loess()` function as an alternative.

## R: Step function

```
library(splines)
m = lm(Sepal.Width~cut(Sepal.Length, 3), data=iris)
summary(m)
```

Call:

```
lm(formula = Sepal.Width ~ cut(Sepal.Length, 3), data = iris)
```

Residuals:

Min	1Q	Median	3Q	Max
-1.2000	-0.2075	-0.0100	0.2000	1.4761

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	3.20000	0.05453	58.686	< 2e-16 ***
cut(Sepal.Length, 3)(5.5,6.7]	-0.27606	0.07378	-3.741	0.000262 ***
cut(Sepal.Length, 3)(6.7,7.9]	-0.09000	0.10837	-0.830	0.407618

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4188 on 147 degrees of freedom

Multiple R-squared: 0.08901, Adjusted R-squared: 0.07662

F-statistic: 7.182 on 2 and 147 DF, p-value: 0.001057



**R: Spline**

```
library(splines)
grids = seq(min(iris$Sepal.Length), max(iris$Sepal.Length), by=0.01)

m = lm(Sepal.Width~bs(Sepal.Length, knots=c(4.3, 5.5, 6.7, 7.9)), data=iris)
summary(m)
```

Call:

```
lm(formula = Sepal.Width ~ bs(Sepal.Length, knots = c(4.3, 5.5,
  6.7, 7.9)), data = iris)
```

Residuals:

```
      Min       1Q   Median       3Q      Max
-1.25514 -0.22417 -0.00521  0.24486  1.40843
```

Coefficients: (2 not defined because of singularities)

	Estimate	Std. Error
(Intercept)	3.26262	0.31055
bs(Sepal.Length, knots = c(4.3, 5.5, 6.7, 7.9))1	-0.50691	0.40431
bs(Sepal.Length, knots = c(4.3, 5.5, 6.7, 7.9))2	0.28136	0.34303
bs(Sepal.Length, knots = c(4.3, 5.5, 6.7, 7.9))3	-0.35769	0.41998
bs(Sepal.Length, knots = c(4.3, 5.5, 6.7, 7.9))4	-0.38126	0.31564
bs(Sepal.Length, knots = c(4.3, 5.5, 6.7, 7.9))5	-0.06646	0.55917
bs(Sepal.Length, knots = c(4.3, 5.5, 6.7, 7.9))6	NA	NA
bs(Sepal.Length, knots = c(4.3, 5.5, 6.7, 7.9))7	NA	NA

	t value	Pr(> t )
(Intercept)	10.506	<2e-16 ***
bs(Sepal.Length, knots = c(4.3, 5.5, 6.7, 7.9))1	-1.254	0.212
bs(Sepal.Length, knots = c(4.3, 5.5, 6.7, 7.9))2	0.820	0.413
bs(Sepal.Length, knots = c(4.3, 5.5, 6.7, 7.9))3	-0.852	0.396
bs(Sepal.Length, knots = c(4.3, 5.5, 6.7, 7.9))4	-1.208	0.229
bs(Sepal.Length, knots = c(4.3, 5.5, 6.7, 7.9))5	-0.119	0.906
bs(Sepal.Length, knots = c(4.3, 5.5, 6.7, 7.9))6	NA	NA
bs(Sepal.Length, knots = c(4.3, 5.5, 6.7, 7.9))7	NA	NA

---


Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.423 on 144 degrees of freedom

Multiple R-squared: 0.08962, Adjusted R-squared: 0.05801

F-statistic: 2.835 on 5 and 144 DF, p-value: 0.01793


```
pred=predict(m,newdata=list(Sepal.Length=grids),se=T)
plot(Sepal.Width~Sepal.Length, data=iris);lines(grids,pred$fit,lwd=2)
```




```
fit=smooth.spline(iris$Sepal.Length,iris$Sepal.Width,cv=T);fit$df
```

```
[1] 16.04326
```

```
plot(Sepal.Width~Sepal.Length, data=iris);lines(fit,col="red",lwd=2)
```



```
m=lm(Sepal.Width~ns(Sepal.Length,df=16),data=iris)
pred=predict(m,newdata=list(Sepal.Length=grids),se=T)
plot(Sepal.Width~Sepal.Length, data=iris);lines(grids, pred$fit,col="red",lwd=2)
```



## Generalized Additive Models

- Allows flexible non-linearities in multiple variables, but retains the additive structure of linear models. For instance, We can fit multiple splines or local regression.
- Use `anova()` to compare models. In the case below, model2 is the best.

```
library(gam)
m1 = gam(Sepal.Width~s(Sepal.Length,df=16), data=iris)
m2 = gam(Sepal.Width~s(Sepal.Length,df=16)+s(Petal.Width,3), data=iris)
m3 = gam(Sepal.Width~s(Sepal.Length,df=16)+s(Petal.Width,3)+Petal.Length, data=iris)
anova(m1,m2,m3,test="F")
```

### Analysis of Deviance Table

```
Model 1: Sepal.Width ~ s(Sepal.Length, df = 16)
Model 2: Sepal.Width ~ s(Sepal.Length, df = 16) + s(Petal.Width, 3)
Model 3: Sepal.Width ~ s(Sepal.Length, df = 16) + s(Petal.Width, 3) +
      Petal.Length
```

	Resid. Df	Resid. Dev	Df	Deviance	F	Pr(>F)
1	133	21.5745				
2	130	10.2482	3.0002	11.3264	52.179	< 2.2e-16 ***
3	129	9.3333	1.0000	0.9149	12.645	0.0005274 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

```
summary(m2)
```

```
Call: gam(formula = Sepal.Width ~ s(Sepal.Length, df = 16) + s(Petal.Width,
      3), data = iris)
```

Deviance Residuals:

	Min	1Q	Median	3Q	Max
	-0.72047	-0.14473	-0.01395	0.17672	0.77893

(Dispersion Parameter for gaussian family taken to be 0.0788)

```
Null Deviance: 28.3069 on 149 degrees of freedom
Residual Deviance: 10.2482 on 130 degrees of freedom
AIC: 65.1512
```

Number of Local Scoring Iterations: 2

### Anova for Parametric Effects

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
s(Sepal.Length, df = 16)	1	0.3696	0.3696	4.6879	0.0322 *
s(Petal.Width, 3)	1	7.0604	7.0604	89.5625	<2e-16 ***
Residuals	130	10.2482	0.0788		

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

### Anova for Nonparametric Effects

	Npar	Df	Npar F	Pr(F)
(Intercept)				
s(Sepal.Length, df = 16)	15	3.023	0.0003461	***
s(Petal.Width, 3)	2	72.446	< 2.2e-16	***


---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1


```
# training MSE
mean((iris$Sepal.Width-predict(m2,newdata=iris))^2)
```

```
[1] 0.06832059
```

```
library(akima)
plot(m2, se=TRUE, col="red")
```



```
grids1 = seq(min(iris$Sepal.Length),max(iris$Sepal.Length), length.out = 50)
grids2 = seq(min(iris$Petal.Length),max(iris$Petal.Length), length.out = 50)
grids3 = seq(min(iris$Petal.Width),max(iris$Petal.Width), length.out = 50)
irisdf = data.frame(Sepal.Length=grids1, Petal.Length=grids2,Petal.Width=grids3)
pred = predict(m2, newdata=irisdf)
plot(Sepal.Width~Sepal.Length, data=iris);lines(grids1, pred,col="red",lwd=2)
```



## Logistic Regression

- Probabilistic, supervised learning, classification
- Linear regression produce a result ranged in  $(-\infty, \infty)$ , which is not appropriate here. Thus, when binary outcome is present, logistic regression is appropriate.
- $\theta(x_i) = \frac{\exp(\beta_0 + \beta_1 x_i)}{1 + \exp(\beta_0 + \beta_1 x_i)} = \frac{1}{1 + \exp(-\{\beta_0 + \beta_1 x_i\})}$
- $\log\left(\frac{\theta(x_i)}{1 - \theta(x_i)}\right) = \beta_0 + \beta_1 x_i$  where  $\log\left(\frac{\theta(x_i)}{1 - \theta(x_i)}\right)$  is called logit.
- $\log\left(\frac{\hat{\theta}(x_i)}{1 - \hat{\theta}(x_i)}\right) = \hat{\beta}_0 + \hat{\beta}_1 x_i$
- We use maximum likelihood estimates for coefficients, but no closed-form solutions exist. Thus, we take advantage of computation results.
- “For every unit increase in X, the odds that the characteristic is present is multiplied by  $\exp(\beta_1)$ .”
- Recall the PDF of binomial:  $P(Y_i = y_i | x_i) = \binom{m_i}{y_i} \theta(x_i)^{y_i} (1 - \theta(x_i))^{m_i - y_i}$
- Recall the log likelihood of binomial :  $\log(L) = \sum_{i=1}^n (y_i(\beta_0 + \beta_1 x_i) - m_i \log(1 + \exp(\beta_0 + \beta_1 x_i))) + \log\left(\binom{m_i}{y_i}\right)$ .

## Goodness of fit test

- $H_o$  : The logistic regression is appropriate vs.  $H_a$  : Not appropriate.
  - We use two log likelihood :  $\log(L_M)$  vs.  $\log(L_S)$  where M refers the logistic regression model and S refers to the saturated model, a model with a theoretically perfect fit. Thus,  $\hat{y}_i = y_i$  under the saturated model.
  - Set  $\hat{y}_i = \hat{\theta}_M(x_i) m_i$
1. Deviance:  $G^2 = 2(\log(L_S) - \log(L_M)) = 2 \sum_{i=1}^n [y_i \log\left(\frac{y_i}{\hat{y}_i}\right) + (m_i - y_i) \log\left(\frac{m_i - y_i}{m_i - \hat{y}_i}\right)] \sim \chi_{n-p-1}^2$ . IF the model is appropriate, then  $G^2$  is smaller so we fail to reject the null. P-value =  $P(\chi_{n-p-1}^2 > G_{obs}^2)$
  2. Pearson  $\chi^2$  statistic:  $\chi^2 = \sum \frac{(y_i/m_i - \hat{\theta}(y_i))^2}{\hat{\theta}(y_i)(1 - \hat{\theta}(y_i))/m_i} \sim \chi_{n-p-1}^2$
  3. R-squared:  $R_{dev}^2 = 1 - \frac{G_{H_a}^2}{G_{H_o}^2}$

## Comparing models

- $H_o : \theta(x) = \frac{1}{1 + \exp(-\beta_0)}$  vs.  $H_a : \theta(x) = \frac{1}{1 + \exp(-\{\beta_0 + \beta_1 x\})}$
- $G_{H_o}^2 - G_{H_a}^2 \sim \chi_{df_1 - df_2}^2$  where  $df_1 = n - (\text{num. of predictors in } H_o) - 1$  and  $df_2 = n - (\text{num. of predictors in } H_a) - 1$

## Marginal model plot

Compare the following two models, and determine wheather the logistic regression is appropriate. If two models are significantly different, then the logistic regression ins not appropriate.

1. Parametric model :  $\theta(x_i) = \frac{1}{1 + \exp(-\{\beta_0 + \beta_1 x_i\})}$
2. Nonparametric model :  $\theta(x) = f(x_1, \dots, x_p)$ . For the model with  $p$  predictors, we need  $p$  many marginal model plots. If any of them show discrepancy from parametric model, then the parametric model (logistic regression) is not appropriate.

## R: Logistic Regression

- Some predictor variables cause an error: `glm.fit: algorithm did not converge` when the sample size is not enough.
- If  $G_{diff} < 0.05$ , reject  $H_0$  and full model is better.

```
m = glm(am~mpg+disp+hp+wt, family=binomial(), data=mtcars)
summary(m)
```

Call:

```
glm(formula = am ~ mpg + disp + hp + wt, family = binomial(),
    data = mtcars)
```

Deviance Residuals:

Min	1Q	Median	3Q	Max
-1.84992	-0.15966	-0.00615	0.01257	1.46081

Coefficients:

	Estimate	Std. Error	z value	Pr(> z )
(Intercept)	-18.48207	40.90451	-0.452	0.651
mpg	1.13503	1.55720	0.729	0.466
disp	-0.02588	0.04087	-0.633	0.527
hp	0.10871	0.09837	1.105	0.269
wt	-4.80560	3.97978	-1.208	0.227

(Dispersion parameter for binomial family taken to be 1)


Null deviance: 43.230 on 31 degrees of freedom  
 Residual deviance: 8.162 on 27 degrees of freedom  
 AIC: 18.162

Number of Fisher Scoring iterations: 9

```
Gdiff = m$null.deviance-m$deviance
pchisq(Gdiff,4,lower=FALSE)
```

```
[1] 4.498811e-07
```

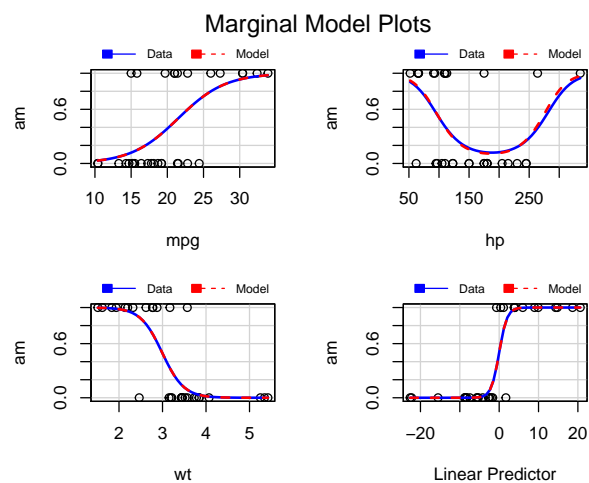
```
library(alr3)
mmps(m)
```



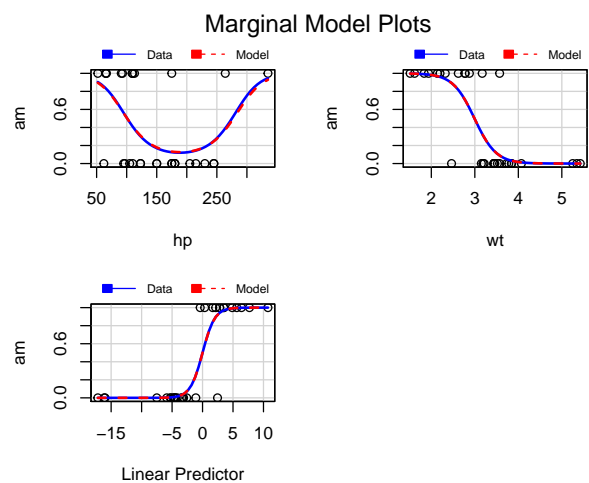
## R: Comparing models

- Here, failed to reject Null, m3 model is better.
- Also, one of the plot shows a discrepancy. Logistic regression for m2 might not be appropriate.

```
m3 = glm(am~mpg+hp+wt, family=binomial(), data=mtcars)
mmps(m3)
```



```
m2 = glm(am~hp+wt, family=binomial(), data=mtcars)
mmps(m2)
```



```
anova(m3,m2,test="Chisq")
```

### Analysis of Deviance Table

Model 1:  $am \sim mpg + hp + wt$

Model 2:  $am \sim hp + wt$

	Resid. Df	Resid. Dev	Df	Deviance	Pr(>Chi)
1	28	8.7661			
2	29	10.0591	-1	-1.293	0.2555

## R: Prediction

```
set.seed(1)
train.i = trainindex(mtcars, trainn = 25)

Y_train = mtcars[train.i, c("am")]
Y_test = mtcars[-train.i, c("am")]
X_train = mtcars[train.i, c("mpg", "hp", "wt")]; X_train = cbind(X_train, Y_train)
X_test = mtcars[-train.i, c("mpg", "hp", "wt")]

m = glm(Y_train ~ mpg + hp + wt, family = binomial(), data = X_train)
Y_pred = ifelse(predict(m, X_test, type = "response") >= 0.5, 1, 0)
table(Y_test, Y_pred)
```

```
      Y_pred
Y_test 0 1
      0 4 0
      1 0 3
```



## Bayes Classifier

- $P() = \frac{P()P()}{P()}$

## K Nearest Neighbor

- Non-probabilistic, supervised learning, classification
- Judge the category of a point via the category of K many Euclidian nearest others.
- No assumption on distribution of X, but they have to be numeric continuous.
- Draw decision boundary according to K many closest neighbors based on Euclidian distance.
- Make sure to standarize the predictor variable X because the algorithm takes Euclidian distance to determine the line, and standarizing increase the accuracy of the results.
- Highly non-linear boundry for smaller K.

---

```
library(class)
set.seed(1)
test.i = sample(1:nrow(Heart),50, replace=F)

# Standarize predictors (continuous)
Xs = scale(cbind(Heart$Age, Heart$MaxHR))
Xs_test = Xs[test.i,]
Xs_train = Xs[-test.i,]
Y_test = Heart$AHD[test.i]
Y_train = Heart$AHD[-test.i]
knn_output = knn(Xs_train, Xs_test, Y_train, k=1)
```

```
# MSE
mean(knn_output!=Y_test)
```

```
[1] 0.4
```

```
# Confusion Matrix
table(knn_output, Y_test)
```

```
      Y_test
knn_output No Yes
No      18   9
Yes     11  12
```

## Find the best K

```
library(caret)
library(e1071)
Heart_train = cbind(Xs_train,Y_train);colnames(Heart_train) = c("Age","MaxHR","AHD")
ctrl = trainControl(method="repeatedcv", repeats=13)
knnfit = train(as.factor(AHD)~., data=Heart_train, trControl=ctrl, method="knn", preProcess = c("center", "scale"))
knnfit
```

k-Nearest Neighbors


253 samples  
 2 predictor  
 2 classes: '1', '2'

Pre-processing: centered (2), scaled (2)  
 Resampling: Cross-Validated (10 fold, repeated 13 times)  
 Summary of sample sizes: 228, 227, 228, 228, 228, 228, ...  
 Resampling results across tuning parameters:

k	Accuracy	Kappa
5	0.6710947	0.3378439
7	0.7098107	0.4156400
9	0.7226105	0.4411433
11	0.7189132	0.4332314
13	0.7192101	0.4333105
15	0.7256154	0.4473893
17	0.7194822	0.4360253
19	0.7118235	0.4208896
21	0.7146055	0.4269985
23	0.7155789	0.4287789

Accuracy was used to select the optimal model using the largest value.  
 The final value used for the model was k = 15.

```
knnfit_val = data.frame(knnfit[4])
plot(knnfit_val[,1], knnfit_val[,2], type="b", col="blue", xlab="k", ylab="Accuracy")
```



```
maxacu = which(knnfit$results$Accuracy==max(knnfit$results$Accuracy));maxacu
```

```
[1] 6
```

## Repeat with the best K

```
knn_output = knn(Xs_train, Xs_test, Y_train, k=maxacu, prob=T)
# MSE
mean(knn_output!=Y_test)
```

```
[1] 0.32
```


```
# Confusion Matrix
table(knn_output, Y_test)
```

	Y_test	
knn_output	No	Yes
No	20	7
Yes	9	14

## Linear Discriminant Analysis

- non-probabilistic, supervised learning, classification.
- Finds the best place to make the best split linear boundary between two (or more) distributions. Relies on the Bayes Classifier. Tries to find the K-dimensional projection that creates the greatest between group separation.
- Assumption: Normality and small sample size, and same variance among all group (category).
- Make the dataset linearly independent. LD1, LD2... are the coefficients such that makes the dataset linearly independent i.e. eigenvectors.
- Dimension reduction.
- The deviation line  $x = \frac{\mu_1 + \mu_2}{2}$

```
ggplot(iris) + geom_density(mapping=aes(x=Sepal.Length, color=as.character(Species))) + geom_vline(xintercept=5.5)
```



- The `lda$svd` represents the eigenvalues, and bigger the better split the data.
- LD1 and LD2 on the bottom represents the trace. Bigger the more important role.
- LD1, LD2 with coefficients are the eigenvectors. Make sure that they are all standardized.

```
library(MASS)
set.seed(1)
test.i = sample(1:nrow(iris), 30, replace=F)
X_train = iris[-test.i, 1:4]
X_test = iris[test.i, 1:4]
Y_train = iris[-test.i, 5]
Y_test = iris[test.i, 5]

# LDA
model_lda = lda(Y_train ~ X_train$Sepal.Length + X_train$Sepal.Width); model_lda
```

Call:

```
lda(Y_train ~ X_train$Sepal.Length + X_train$Sepal.Width)
```

Prior probabilities of groups:

setosa	versicolor	virginica
0.3166667	0.3416667	0.3416667

Group means:

```

      X_train$Sepal.Length X_train$Sepal.Width
setosa           5.018421           3.418421
versicolor       5.926829           2.760976
virginica        6.560976           2.960976

```

Coefficients of linear discriminants:

```

              LD1      LD2
X_train$Sepal.Length -2.074508 -0.8362372
X_train$Sepal.Width  2.846646 -2.0806011

```

Proportion of trace:

```

      LD1      LD2
0.9602 0.0398


```

```

# Take LD1(The first split)
LD1 = predict(model_lda)$x[,1]

plot(Y_test)


```



```

plot(LD1, type="n");text(LD1,labels=unclass(iris$Species))

```



```
m3 = lda(Species ~ Sepal.Length + Sepal.Width + Petal.Length + Petal.Width, data = iris)
m3
```

Call:

```
lda(Species ~ Sepal.Length + Sepal.Width + Petal.Length + Petal.Width,
    data = iris)
```

Prior probabilities of groups:

	setosa	versicolor	virginica
	0.3333333	0.3333333	0.3333333

Group means:

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
setosa	5.006	3.428	1.462	0.246
versicolor	5.936	2.770	4.260	1.326
virginica	6.588	2.974	5.552	2.026


Coefficients of linear discriminants:

	LD1	LD2
Sepal.Length	0.8293776	0.02410215
Sepal.Width	1.5344731	2.16452123
Petal.Length	-2.2012117	-0.93192121
Petal.Width	-2.8104603	2.83918785

Proportion of trace:

	LD1	LD2
	0.9912	0.0088

```
LD1<-predict(m3)$x[,1]
LD2<-predict(m3)$x[,2]
plot(LD1,LD2,xlab="first linear discriminant",ylab="second linear discriminant",type="n")
text(cbind(LD1,LD2),labels=unclass(iris$Species))
```



```
head(2.105107+0.8293776*iris$Sepal.Length+1.5344731*iris$Sepal.Width-2.2012117*iris$Petal.Length-2.8104603*iris$Petal.Width)
```

```
[1] 8.061800 7.128688 7.489828 6.813201 8.132310 7.701947
```

```
head(LD1)
```

	1	2	3	4	5	6
	8.061800	7.128688	7.489828	6.813201	8.132309	7.701947

```
cor(iris[,1],LD1)
```

```
[1] -0.7918878
```

```
m3$svd
```

```
[1] 48.642644 4.579983
```

```
iris.lda<-lda(Species ~ Sepal.Length + Sepal.Width + Petal.Length + Petal.Width, data = iris)
iris.lda
```

Call:

```
lda(Species ~ Sepal.Length + Sepal.Width + Petal.Length + Petal.Width,
    data = iris)
```

Prior probabilities of groups:

```
      setosa versicolor virginica
0.3333333 0.3333333 0.3333333
```

Group means:

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
setosa	5.006	3.428	1.462	0.246
versicolor	5.936	2.770	4.260	1.326
virginica	6.588	2.974	5.552	2.026


Coefficients of linear discriminants:

	LD1	LD2
Sepal.Length	0.8293776	0.02410215
Sepal.Width	1.5344731	2.16452123
Petal.Length	-2.2012117	-0.93192121
Petal.Width	-2.8104603	2.83918785

Proportion of trace:

	LD1	LD2
	0.9912	0.0088

```
plot(iris.lda)
```



```
iris.lda$svd
```

```
[1] 48.642644 4.579983
```

```
iris.lda$counts
```



```

      setosa versicolor virginica
      50         50         50

iris.lda$means

      Sepal.Length Sepal.Width Petal.Length Petal.Width
setosa           5.006      3.428         1.462         0.246
versicolor       5.936      2.770         4.260         1.326
virginica        6.588      2.974         5.552         2.026

iris.lda$lev

[1] "setosa"      "versicolor"  "virginica"

```


```

# Plots:
LD1<-predict(iris.lda)$x[,1]
LD2<-predict(iris.lda)$x[,2]

plot(LD1,LD2,xlab="first linear discriminant",ylab="second linear discriminant",col=2:4,type="n",main="")
text(cbind(LD1,LD2),labels=iris$Species)
legend(0.5,2.8,legend=c("Setosa","Versicolor","Virginica"))

```

LDA1 vs LDA2 for the three Species




```

plot(LD1,LD2,xlab="first linear discriminant",ylab="second linear discriminant",col=2:4,type="n",main="")
text(cbind(LD1,LD2),labels=unclass(iris$Species))
legend(0.5,2.8,legend=c("1=Setosa","2=Versicolor","3=Virginica"))


```

LDA1 vs LDA2 for the three Species



```
# 1="setosa"
# 2="versicolor"
# 3="virginica"
```

```
qplot(x = LD1, y = LD2, colour = iris$Species, shape = iris$Species, main="LDA1 vs LDA2 for the three Species")
```



```
# Group centroids
```

```
sum(LD1*(iris$Species=="setosa"))/sum(iris$Species=="setosa")
```

```
[1] 7.6076
```

```
sum(LD2*(iris$Species=="setosa"))/sum(iris$Species=="setosa")
```

```
[1] 0.215133
```

```
sum(LD1*(iris$Species=="versicolor"))/sum(iris$Species=="versicolor")
```

```
[1] -1.825049
```

```
sum(LD2*(iris$Species=="versicolor"))/sum(iris$Species=="versicolor")
```

```
[1] -0.7278996
```

```
sum(LD1*(iris$Species=="virginica"))/sum(iris$Species=="virginica")
```

```
[1] -5.78255
```

```
sum(LD2*(iris$Species=="virginica"))/sum(iris$Species=="virginica")
```

```
[1] 0.5127666
```

```
iris.predict<-predict(iris.lda,iris[,1:4])
```

```
iris.classify<-iris.predict$class
```

```
iris.classperc<-sum(iris.classify==iris[,5])/150
```

```
iris.classperc
```

```
[1] 0.98
```

```
table(Original=iris$Species,Predicted=predict(iris.lda)$class)
```

Original	Predicted		
	setosa	versicolor	virginica
setosa	50	0	0
versicolor	0	48	2
virginica	0	1	49

## Quadratic Discriminant Analysis

- In stead of linear boundary, we use non-linear boundary.

```
m4 = qda(Species ~ Sepal.Length + Sepal.Width + Petal.Length + Petal.Width, data = iris)
m4
```

Call:

```
qda(Species ~ Sepal.Length + Sepal.Width + Petal.Length + Petal.Width,
    data = iris)
```

Prior probabilities of groups:

setosa	versicolor	virginica
0.3333333	0.3333333	0.3333333

Group means:

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
setosa	5.006	3.428	1.462	0.246
versicolor	5.936	2.770	4.260	1.326
virginica	6.588	2.974	5.552	2.026

## Support Vector Machine

## K-Means Clustering

## Kernelized Clustering

---

## EM Type Algorithm

## Decision Trees


- Partition data points. Determine the value of response variable (if continuous, the mean of the response) according to which partition that new predictor belongs to (looks like a house diagram). If Y is continuous, it is called regression tree and if Y is categorical, it is called classification tree.
- Determine the splits lines that generate the lowest MSE. As we have more splits, the new value of Y can be determined on the tree structure, or “nested if” structure.
- Pruning tree: A large tree can be over-fitting, thus pruning tree allows cutting off some of the terminal nodes. Find the best pruning by cross validation.
- If the relationship between the predictors and response is linear, then linear regression is better; if the relationship is “rectangle” shapes, tree performs better.
- Easily interpretable (e.g. If weight is above 40kg, height is higher than 140cm in average). However, suffer from high variance. Thus, use random forest to resolve this issue.

## R: Tree

```
library(tree)
train.i = trainindex(iris)
iris_train = iris[train.i,]
iris_test = iris[-train.i,1:4]
Y_test = factor(iris[-train.i,5])

tree1 = tree(Species~., data=iris_train)
summary(tree1)
```

```
Classification tree:
tree(formula = Species ~ ., data = iris_train)
Variables actually used in tree construction:
[1] "Petal.Length" "Petal.Width"  "Sepal.Length"
Number of terminal nodes: 6
Residual mean deviance: 0.1329 = 17.14 / 129
Misclassification error rate: 0.02963 = 4 / 135
plot(tree1);text(tree1)
```





```
# Prediction
preds=predict(tree1,newdata=iris_test, type="class")
table(Y_test, preds)
```


Y_test	preds		
	setosa	versicolor	virginica
setosa	5	0	0
versicolor	0	6	0
virginica	0	0	4

---


## R: Pruning Tree

- Here, it shows that pruning with 5 branches is appropriate.

```
cv.train=cv.tree(tree1,FUN=prune.misclass)
plot(cv.train$dev~cv.train$size)
```



```
pruned.fit=prune.misclass(tree1,best=5)
plot(pruned.fit)
text(pruned.fit,pretty=TRUE)
```



```
summary(pruned.fit)
```

Classification tree:

```
tree(formula = Species ~ ., data = iris_train)
```

Variables actually used in tree construction:

```
[1] "Petal.Length" "Petal.Width" "Sepal.Length"
```

Number of terminal nodes: 6

Residual mean deviance: 0.1329 = 17.14 / 129

Misclassification error rate: 0.02963 = 4 / 135

```
# Prediction
```

```
pred.prune=predict(pruned.fit,newdata=iris_test,type="class")
```

```
table(Y_test, pred.prune)
```

	pred.prune		
Y_test	setosa	versicolor	virginica
setosa	5	0	0
versicolor	0	6	0
virginica	0	0	4

---

## R: Regression Tree

```
mtcars_temp = subset(mtcars,select=c(mpg,wt, disp, qsec))
```

```
tree2=tree(mpg~., data=mtcars_temp)
```

```
summary(tree2)
```

Regression tree:

```
tree(formula = mpg ~ ., data = mtcars_temp)
```

Number of terminal nodes: 5


Residual mean deviance: 5.104 = 137.8 / 27

Distribution of residuals:

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
-4.067	-1.637	0.100	0.000	1.338	3.833

```
plot(tree2)
```

```
text(tree2,pretty=0)
```



## Random Forest

- Use bagging: Bootstrap + averaging. That is, generate B different bootstrapped training dataset. Train the statistical learning method on each of the B training datasets, and obtain the prediction, then take the average. In this case, construct B different trees using B bootstrapped dataset, then take the average of the results.
- If continuous, average all predictions from all B trees. If Classification, majority vote among all B trees. These trees are not pruned, so each individual tree has high variance but low bias. Averaging these trees reduces variance, and thus lowering both variance and bias can be achieved.
- Two methods for prediction: Record the class that each bootstrapped data set predicts and provide an overall prediction to the most commonly occurring one (majority vote). Or, if our classifier produces probability estimates we can just average the probabilities and then predict to the class with the highest probability.
- Bagging improves prediction accuracy at the expense of interpretability. But, we can use relative influence plots to see the contributions of each variables to the model. Larger the more influential.
- Random forest: Build a number of decision trees on bootstrapped training sample, but when building these trees, each time a split in a tree is considered, a random sample of m predictors is chosen as split candidates from the full set of p predictors. Only m predictors are used for the sake of “de-correlation” of the model; if all variables are used, each models for each bootstrapped dataset will be similar to each other, hence highly correlated. Averaging many highly correlated quantities does not lead to a large variance reduction.

---


### R: RF, mtry=4

```
library(randomForest)
```

```
RF=randomForest(Species~.,data=iris_train,mtry=4,importance=TRUE, ntree=100)
summary(RF)
```

	Length	Class	Mode
call	6	-none-	call
type	1	-none-	character
predicted	135	factor	numeric
err.rate	400	-none-	numeric
confusion	12	-none-	numeric
votes	405	matrix	numeric
oob.times	135	-none-	numeric
classes	3	-none-	character
importance	20	-none-	numeric
importanceSD	16	-none-	numeric
localImportance	0	-none-	NULL
proximity	0	-none-	NULL
ntree	1	-none-	numeric
mtry	1	-none-	numeric
forest	14	-none-	list
y	135	factor	numeric
test	0	-none-	NULL
inbag	0	-none-	NULL
terms	3	terms	call

```
plot(RF)
```



```
print(RF)
```

Call:

```
randomForest(formula = Species ~ ., data = iris_train, mtry = 4, importance = TRUE, ntree = 100)
      Type of random forest: classification
      Number of trees: 100
No. of variables tried at each split: 4
```

OOB estimate of error rate: 5.93%

Confusion matrix:

	setosa	versicolor	virginica	class.error
setosa	45	0	0	0.00000000
versicolor	0	41	3	0.06818182
virginica	0	5	41	0.10869565

```
importance(RF)
```

	setosa	versicolor	virginica	MeanDecreaseAccuracy
Sepal.Length	0.000000	3.534282	-0.1174515	3.0153919
Sepal.Width	0.000000	-1.549744	1.7981879	0.3076241
Petal.Length	10.658224	16.926435	13.5961096	16.4290251
Petal.Width	9.700194	15.143134	11.7300800	14.0252231

	MeanDecreaseGini
Sepal.Length	1.150223
Sepal.Width	1.013324
Petal.Length	41.384468
Petal.Width	45.788132

```
varImpPlot (RF)
```




Y_test	preds		
	setosa	versicolor	virginica
setosa	5	0	0
versicolor	0	6	0
virginica	0	0	4

R: RF, mtry=2

```
RF=randomForest(Species~.,data=iris_train,mtry=2,importance=TRUE, ntree=100)
summary(RF)
```

	Length	Class	Mode
call	6	-none-	call
type	1	-none-	character
predicted	135	factor	numeric
err.rate	400	-none-	numeric
confusion	12	-none-	numeric
votes	405	matrix	numeric
oob.times	135	-none-	numeric
classes	3	-none-	character
importance	20	-none-	numeric
importanceSD	16	-none-	numeric
localImportance	0	-none-	NULL
proximity	0	-none-	NULL
ntree	1	-none-	numeric
mtry	1	-none-	numeric
forest	14	-none-	list
y	135	factor	numeric
test	0	-none-	NULL
inbag	0	-none-	NULL
terms	3	terms	call

```
plot(RF)
```



```
print(RF)
```

Call:  
randomForest(formula = Species ~ ., data = iris\_train, mtry = 2, importance = TRUE, ntree = 100)  
Type of random forest: classification  
Number of trees: 100  
No. of variables tried at each split: 2

OOB estimate of error rate: 5.19%  
Confusion matrix:  
          setosa versicolor virginica class.error  
setosa      45          0          0 0.00000000  
versicolor   0         41          3 0.06818182  
virginica    0          4         42 0.08695652

```
importance(RF)
```


	setosa	versicolor	virginica	MeanDecreaseAccuracy
Sepal.Length	3.003956	4.01824632	1.885682	4.663758
Sepal.Width	2.015281	0.02670148	2.832304	2.617572
Petal.Length	11.846486	15.56040968	13.867746	16.436637
Petal.Width	8.171265	14.30018424	13.637917	13.928544

	MeanDecreaseGini
Sepal.Length	9.618752
Sepal.Width	2.143676
Petal.Length	42.021088
Petal.Width	35.507632

```
varImpPlot (RF)
```

RF



Y_test	preds		
	setosa	versicolor	virginica
setosa	5	0	0
versicolor	0	6	0
virginica	0	1	3

## Neural Network

- Use CV to adjust the size of nodes.


```
library(nnet)
library(NeuralNetTools)

train.i = trainindex(iris)
iris_train = iris[train.i,]
iris_test = iris[-train.i,1:4]
Y_test = iris[-train.i,5]

m = nnet(Species ~ Sepal.Length + Sepal.Width + Petal.Length + Petal.Width, size=3, data = iris_train)

# weights: 27
initial value 162.430958
iter 10 value 59.768659
iter 20 value 7.339820
iter 30 value 6.328510
iter 40 value 5.930743
iter 50 value 4.868825
iter 60 value 4.773977
iter 70 value 4.773034
iter 80 value 4.772938
final value 4.772928
converged

plotnet(m)
```



```
results = max.col(predict(m, iris_test))
table(Y_test, results)
```

```
      results
Y_test  1 2 3
setosa   4 0 0
versicolor 0 4 0
virginica 0 0 7
```



## Comparing all classification methods

Method	Linearity	Normality	Constant Variance	Big Sample Size	K>2
Logistic Regression	Yes	Yes on Y	Yes	Yes	K=2
LDA	Yes	Yes on X and Y	Yes	No	Yes
QDA	No	Yes	No	No	Yes
KNN	No	No	No	Yes	Yes

## Experimental Design

- Enable to study cause and effect relationship.
  - Three key factors: Randomization, Replication, Blocking.
1. Randomization: The allocation of units to treatment must be randomly determined to prevent subjective assignments and possible biases. The order of the experiment also must be randomized. Also, “average out” the effects irrelevant or unknown factors that might be present.
  2. Replication: Each treatment is applied to a different experimental unit. For each treatment, good number of replications are desired to obtain more accurate estimate of experimental error and results. Repetition refers to measuring the same observation with the same factors over multiple times.
  3. Blocking: Dealing with nuisance factors, the factors that may influence the experimental results but not interested and cannot be controlled. If there is a factor that can be controlled e.g. gender, block them. If not, rely on randomization.
- Guideline
    1. Recognition of and statement of problem
    2. Selection of response variable, continuous response is preferable.
    3. Choice of factors, levels, ranges.
    4. Choice of experimental design
    5. Perform experiment
    6. Analyze the data
    7. Conclusion/recommendation
  - Assumptions are the same as OLS: Independence, constant variance, normality.
  - Type I error:  $\alpha = P(\text{reject Null} | \text{Null is true})$
  - Type II error:  $\beta = P(\text{fail to reject Null} | \text{Null is false})$

### t-test

- Hypothesis:  $H_0 : \mu_1 = \mu_2 = \dots = \mu_k$
- $t_o = \frac{\bar{y}_1 - \bar{y}_2}{S_p \sqrt{1/n_1 + 1/n_2}} \sim t_{n_1 + n_2 - 2}$


## Completely Randomized Design

- Key: one treatment with many levels.
- If no blocking, its structure and test style is very similar (or identical) to two sample t-test.
- Randomized treatment and blocking order to prevent biases. The role of repetition is to reduce the variation on the result of the experiment.
- Effect model  $y_{ij} = \mu + \tau_i + \epsilon_{ij}$  for  $i = 1, \dots, a$ ,  $j = 1, \dots, n$ .

### R: No Blocking Case

```
bond = c(16.85,16.40,17.21,16.35,16.52,17.04,16.96,17.15,16.59,16.57,16.62,16.75,17.37,17.12,16.98,16.85)
mortar = c(rep("modified",10),rep("unmodified",10))
```

```
# Check normality assumption
par(mfrow=c(1,2))
qqnorm(bond[mortar=="modified"]);qqline(bond[mortar=="modified"])
qqnorm(bond[mortar=="unmodified"]);qqline(bond[mortar=="unmodified"])
```



```
par(mfrow=c(1,1))
# Assuming unequal variances:
# numeric ~ categorical
m1=t.test(bond~mortar, var.equal=FALSE, alternative="two.sided", mu=0);m1
```

#### Welch Two Sample t-test

```
data: bond by mortar
t = -2.1869, df = 17.025, p-value = 0.043
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 -0.546174139 -0.009825861
sample estimates:
 mean in group modified mean in group unmodified
          16.764          17.042
```

```
m1$p.value
```

```
[1] 0.04299838
```

```
# Equal Var
```

```
m2=t.test(bond~mortar, var.equal=TRUE);m2
```

Two Sample t-test

data: bond by mortar

t = -2.1869, df = 18, p-value = 0.0422

alternative hypothesis: true difference in means is not equal to 0

95 percent confidence interval:

-0.54507339 -0.01092661

sample estimates:

mean in group modified mean in group unmodified

16.764

17.042

```
m2$p.value
```

```
[1] 0.04219672
```

## Paired Comparison Design

- Key: one treatment with two levels only and one block.
- Special case of RBD. Make comparisons within matched pairs of experimental material.
- Effect model  $y_{ij} = \mu_i + \beta_j + \epsilon_{ij}$  for  $i = 1, 2, j = 1, \dots, n$ .
- $H_o : \mu_d = \mu_1 - \mu_2 = 0$  and  $t = \frac{\bar{d}}{S_d/\sqrt{n}} \sim t_{n-1}$  OR  $H_o : \sigma_1^2 = \sigma_2^2$  and  $F_0 = \frac{S_1^2}{S_2^2} \sim F_{n_1-1, n_2-1}$
- Under this synario, blocking reduces variance much more than CRBD, because its degrees of freedom is much bigger than this design. In words, df determines the sensitivity of the test: bigger the more sensitive and wider CI.
- Blocking is not always the best design strategy. If the within block variability is the same as the between block variability, the variance of will be the same regardless of which design is used. Actually, blocking in this situation would be a poor choice of design because blocking results in the loss of n-1 degrees of freedom and will actually lead to a wider confidence interval on  $\mu_1 - \mu_2$ .

	CRD	RBD
Null	$\mu_1 - \mu_2 = 0$	$\mu_d = 0$
Sample Stats	$\bar{y}_1 - \bar{y}_2$	$\bar{d}$
Test Stats	$t_o = \frac{\bar{y}_1 - \bar{y}_2}{S_1 \sqrt{1/n_1 + 1/n_2}}$	$t = \frac{\bar{d}}{S_d/\sqrt{n}}$
Est. of SD	$S_p = \sqrt{\frac{(n_1-1)S_1^2 + (n_2-1)S_2^2}{n_1 + n_2 - 2}}$	$S_d = \sqrt{\frac{\sum (d_j - \bar{d})^2}{n-1}}$
DF	$n_1 + n_2 - 2$	$n - 1$

## R: Paired Deisgn

```
hardness <- c(7,3,3,4,8,3,2,9,5,4,6,3,5,3,8,2,4,9,4,5)
tip <- c(rep("tip1",10),rep("tip2",10))
di <- hardness[tip=="tip1"]-hardness[tip=="tip2"]
data2 <- data.frame(hardness=hardness,tip=tip, di=di)
head(data2)
```

```
hardness tip di
1         7 tip1 1
2         3 tip1 0
3         3 tip1 -2
4         4 tip1 1
5         8 tip1 0
6         3 tip1 1
```

```
n <- 10
se_paired <- sd(di)/sqrt(n)
se_paired
```

```
[1] 0.3785939
```

```
t.test(hardness~tip, paired=TRUE)
```

### Paired t-test

```
data: hardness by tip
t = -0.26414, df = 9, p-value = 0.7976
```

```

alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 -0.9564389  0.7564389
sample estimates:
mean of the differences
          -0.1

```

---

## R: Test for equal variance

```
var.test(bond~mortar, alternative="two.sided")
```

F test to compare two variances

```

data:  bond by mortar
F = 1.6293, num df = 9, denom df = 9, p-value = 0.4785
alternative hypothesis: true ratio of variances is not equal to 1
95 percent confidence interval:
 0.4046845 6.5593806
sample estimates:
ratio of variances
      1.629257

# p-value > 0.05, fail to reject the null : can assume the equal variance.
# Appropriate t-test setting :
t.test(bond~mortar, var.equal=TRUE)

```

Two Sample t-test

```

data:  bond by mortar
t = -2.1869, df = 18, p-value = 0.0422
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 -0.54507339 -0.01092661
sample estimates:
mean in group modified mean in group unmodified
          16.764              17.042

```

## ANOVA for experiment

### Basics

- One-way ANOVA for one factor i.e. slope of regression, and Two-way ANOVA for two-factor i.e. treatment and controll. Use F-test.
- Means Model  $y_{ij} = \mu_i + \epsilon_{ij}$  for  $i = 1, \dots, a, j = 1, \dots, n$
- Effect model  $y_{ij} = \mu + \tau_i + \epsilon_{ij}$  for  $i = 1, \dots, a, j = 1, \dots, n$ .
- Hypothesis:  $H_o : \mu_1 = \mu_3 = \dots = \mu_a$

	SS	DF	MS	F
Between treatments	$SS_{treat} = n \sum_{i=1}^a (\bar{y}_i - \bar{\bar{y}})^2$	a-1	$MS_{treat} = \frac{SS_{treat}}{n-1}$	$F = \frac{MS_{treat}}{MS_E}$
Error	$SS_E = SS_T - SS_{treat}$	N-a	$MS_E$	
Total	$SS_T = \sum_{i=1}^a \sum_{j=1}^n (y_{ij} - \bar{\bar{y}})^2$	N-1		

```
require(Hmisc)
Diet = c("Control","Control","Sucrose","Sucrose", "Glucose","Glucose","Fructose","Fructose")
Times = c(2.3,1.7,4.0,3.6,2.9,2.7,2.1,2.3)
summary(Times~Diet)
```

Times      N= 8

```
+-----+-----+-----+
|      |      |N|Times|
+-----+-----+-----+
|Diet  |Control|2|2.0  |
|      |Fructose|2|2.2  |
|      |Glucose|2|2.8  |
|      |Sucrose|2|3.8  |
+-----+-----+-----+
|Overall|      |8|2.7  |
+-----+-----+-----+
```

## R: Effect Model

- Result:  $\tau_1 = -0.7$ ,  $\tau_2 = -0.5$  and so forth. Thus, the effect of Control is 0.7 lower than the ground mean. Diet is significant, so we reject the null and conclude that at least one treatment is significant i.e. mean is different.

```
m = aov(Times~Diet)
model.tables(m)
```

Tables of effects

```
Diet
Diet
Control Fructose Glucose Sucrose
-0.7      -0.5      0.1      1.1
```

```
summary(m)
```

```
          Df Sum Sq Mean Sq F value    Pr(>F)
Diet         3   3.92   1.307    17.42 0.00925 **
Residuals    4   0.30   0.075
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

## R: Means Model

- Null for ANOVA is rejected, so at least one of the treatment is significant. In fact, we know that Glucose and Sucrose are.

```
m = lm(Times~Diet)
summary(m)
```

Call:

```
lm(formula = Times ~ Diet)
```

Residuals:

```
 1    2    3    4    5    6    7    8
0.3 -0.3  0.2 -0.2  0.1 -0.1 -0.1  0.1
```

Coefficients:

```
              Estimate Std. Error t value Pr(>|t|)
(Intercept)   2.0000     0.1936  10.328 0.000496 ***
DietFructose   0.2000     0.2739   0.730 0.505681
DietGlucose    0.8000     0.2739   2.921 0.043192 *
DietSucrose    1.8000     0.2739   6.573 0.002773 **
```

---

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Residual standard error: 0.2739 on 4 degrees of freedom

Multiple R-squared: 0.9289, Adjusted R-squared: 0.8756

F-statistic: 17.42 on 3 and 4 DF, p-value: 0.009248

```
anova(m)
```

Analysis of Variance Table



Response: Times

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
Diet	3	3.92	1.3067	17.422	0.009248 **
Residuals	4	0.30	0.0750		

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

## Model Adequacy

- Residuals:  $e_{ij} = y_{ij} - \bar{y}_{i.}$ . Check the following:
  1. Residuals are normally distributed (use normal probability plot)
  2. the residuals are random and independent one another
  3. the variance of the residuals is constant over different treatments. Use residual vs. fitted plot

## Post-hoc analysis

- Examine all pairs(i,j) to find which specific group means are significantly different one another. Confidence interval for  $\mu_i - \mu_j$  for all pairs. Hypothesis testing for  $H_0 : \mu_i = \mu_j$  for all pairs.

-Fisher's LSD Method

- Tukey's Method.

## Power of the test

- Obtain number of replications. Use power:  $1 - \beta$ , the probability of correctly rejecting Null given the Null is false. In ANOVA, we conclude the treatment means are significantly different when the true means are really different.
- 1. Power depends on the significant level  $\alpha$ . If you reduce the significance level (e.g., from 0.05 to 0.01), the rejection region gets smaller. As a result, you are less likely to reject the null hypothesis. This means that you are less likely to reject the null hypothesis when it is false, so you are more likely to make a Type II error. In short, the power of the test is reduced when you reduce the significance level ; and vice versa.
- 2. Sample size n. A larger sample size narrows the distribution of the test statistic. Thus, there is a less overlap between the sampling distribution under  $H_0$  vs  $H_1$ . For the fixed significance level, power will increase as the sample size increases. In fact, BOTH Power and  $1 - \alpha$  increase, because as n increases, we are more likely to make correct decision.
- 3. Effect size (f). To know if an observed difference is not only statistically significant but also important or meaningful, you will need to calculate its effect size:  $f = \sqrt{\frac{\sum (\mu_i - \mu)^2}{a\sigma^2}}$ , which measures how far each treatment mean must be from the grand mean, but it is not common to know the true means. We usually use the meaningful difference we would be interested in: the maximum difference among each treatment means  $d = \max(\mu_i - \mu_j)$ .

-`pwr.anova.test(k,n.f.sig.level=0.05,power)`. K = number of treatment. Need three input, then R will return the missing forth one.

---

```
library(pwr)
pwr.anova.test(k=3,f=0.59, power=0.9)
```

Balanced one-way analysis of variance power calculation

```
k = 3
n = 13.17395
f = 0.59
sig.level = 0.05
power = 0.9
```

NOTE: n is number in each group

```
pwr.anova.test(k=3,n=10,f=0.3)
```

Balanced one-way analysis of variance power calculation

```
k = 3
n = 10
f = 0.3
sig.level = 0.05
power = 0.2655592
```

NOTE: n is number in each group

## Completely Randomized Block Design

- Key: One treatment, one block with many levels.
- In CRD, we did not consider nuisance factor. Here, we do. blocking is one approach, especially if the nuisance factors are known and controllable.
- This design strategy improves the accuracy of the comparisons among treatments by eliminating the variability among the blocks.
- Generalization of paired comparison design.
- $y_{ij} = \mu + \tau_i + \beta_j + \epsilon_{ij}$  for  $i = 1, \dots, a$ ,  $j = 1, \dots, b$  where  $\mu$  is overall mean,  $\tau_i$  is the effect of the  $i$ th treatment,  $\beta_j$  is the effect of  $j$ th factor.
- $H_o : \tau_1 = \dots = \tau_a = 0$ , use F test where  $F_0 \sim F_{a-1, (a-1)(b-1)}$
- To test if blocking is necessary,  $H_o : \beta_1 = \dots = \beta_b = 0$ . However,  $F_0 = \frac{MS_{blocks}}{MS_E}$  CANNOT be used.

	SS	DF	MS	F
Treatments	$SS_{treat}$	a-1	$MS_{treat} = \frac{SS_{treat}}{a-1}$	$F = \frac{MS_{treat}}{MS_E}$
Blocks	$SS_{block}$	b-1	$MS_{block} = \frac{SS_{block}}{b-1}$	
Error	$SS_E =$ $SS_T - SS_{treat} - SS_{block}$	(a-1)(b-1)	$MS_E = \frac{SS_E}{(a-1)(b-1)}$	
Total	$SS_T$	N-1		

## Latin Square Design

- Key: one treatment, two blocks with all same levels.
- Latin Squares design deals with the case where there are two nuisance factors to control.
- Both nuisance factors and treatment must have the same number of levels. Only one observation per combination. Latin square designs are reasonable choices when it is impossible to use each treatment level for the same combination of blocking levels. The design is a square arrangement and that p treatments are denoted by the Latin letters A, B, C, . . . ; hence the name Latin square.
- 1 treatments (p levels), 2 nuisance (p levels each)
- $y_{ijk} = \mu + \alpha_i + \tau_j + \beta_k + \epsilon_{ij}$  for  $i = 1, \dots, p, j = 1, \dots, p$  and  $k = 1, \dots, p$ , where  $\mu$  is overall mean,  $\tau_j$  is the effect of the jth treatment,  $\alpha_i$  is the row effect,  $\beta_k$  is the column effect.
- $H_o : \tau_1 = \dots = \tau_p = 0$ , use F test where  $F_0 F_{p-1, (p-2)(p-1)}$

## Greco-Latin Square Design

- Key: one treatment, three blocks with all same levels.
- Latin Square Design + One more nuisance factor with p levels.
- $y_{ijkl} = \mu + \theta_i + \tau_j + \omega_k + \psi_l + \epsilon_{ijk}$  where  $i, j, k, l = 1 \dots p$ .
- $H_o : \tau_1 = \dots = \tau_p = 0$ , use F test where  $F_0 F_{p-1, (p-3)(p-1)}$

## R: Generate Latin Square Design Structure

```
library(agricolae)
T1<-c("A", "B", "C", "D")
T2<-c("a", "b", "c", "d")
design.graeco(T1,T2)$sketch
```

```
      [,1] [,2] [,3] [,4]
[1,] "B d" "D a" "C c" "A b"
[2,] "D c" "B b" "A d" "C a"
[3,] "C b" "A c" "B a" "D d"
[4,] "A a" "C d" "D b" "B c"
```

## R: ANOVA and testing

```
row <- c(rep(1,5),rep(2,5),rep(3,5),rep(4,5),rep(5,5))
column <-c(rep(1:5,5))
Latin <- c("A","B","C","D","E","B","C","D","E","A","C","D","E","A","B","D","E","A","B","C","E","A","B",
Greek <- c("a","r","e","b","d","b","d","a","r","e","r","e","b","d","a","d","a","r","e","b","e","b","d",
burning <- c(-1,-5,-6,-1,-1,-8,-1,5,2,11,-7,13,1,2,-4,1,6,1,-2,-3,-3,5,-5,4,6)
data.frame(row=row,column=column, Latin=Latin, Greek=Greek, burning=burning)
```

```
      row column Latin Greek burning
1      1      1      A      a      -1
2      1      2      B      r      -5
3      1      3      C      e      -6
```

4	1	4	D	b	-1
5	1	5	E	d	-1
6	2	1	B	b	-8
7	2	2	C	d	-1
8	2	3	D	a	5
9	2	4	E	r	2
10	2	5	A	e	11
11	3	1	C	r	-7
12	3	2	D	e	13
13	3	3	E	b	1
14	3	4	A	d	2
15	3	5	B	a	-4
16	4	1	D	d	1
17	4	2	E	a	6
18	4	3	A	r	1
19	4	4	B	e	-2
20	4	5	C	b	-3
21	5	1	E	e	-3
22	5	2	A	b	5
23	5	3	B	d	-5
24	5	4	C	a	4
25	5	5	D	r	6

### R: Latin Square

- Reject Null for the treatment

```
latinmodel <- aov(burning~factor(Latin)+factor(row)+factor(column))
summary(latinmodel)
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
factor(Latin)	4	330	82.50	7.734	0.00254 **
factor(row)	4	68	17.00	1.594	0.23906
factor(column)	4	150	37.50	3.516	0.04037 *
Residuals	12	128	10.67		

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

### R: Greco-Latin Square

- Reject Null for the treatment

```
glmodel <- aov(burning~factor(Latin)+factor(row)+factor(column)+factor(Greek))
summary(glmodel)
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
factor(Latin)	4	330	82.50	10.000	0.00334 **
factor(row)	4	68	17.00	2.061	0.17831
factor(column)	4	150	37.50	4.545	0.03293 *
factor(Greek)	4	62	15.50	1.879	0.20764
Residuals	8	66	8.25		

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

## Balanced Incomplete Design

- Key: Cannot run all combinations, one treatment, one block.
  - Sometimes we may not be able to run all the treatment combinations in each block for randomized block design experiments. A balanced incomplete block design (BIBD) is an incomplete block design in which any two treatments appear together an equal number of times.
  - “a” treatments and “b” blocks. Each block contains “k” treatments that each treatment occurs “r” times. Thus,  $N = ar = bk$ .  $k < a$ . Each pair of treatments occurs together the same number of times in total (Check this via  $\lambda = \frac{r(k-1)}{a-1}$  that must be an integer.)
  - $y_{ij} = \mu + \tau_i + \beta_j + \epsilon_{ij}$  where  $i = 1 \dots a$ ,  $j = 1 \dots b$ .
  - Again, use F test (search ANOVA online, it’s easier...)
- 

### R: BIBD structure

- Num. of treatment  $a = 4$
- Num. of block  $b = 6$
- Num. of treatment / block  $k = 2$
- Num. of replication / treatment  $r = 3$
- $\lambda = 1$

```
require(crossdes)
set.seed(1)
find.BIB(4, 6, 2)
```

```
      [,1] [,2]
[1,]     2     3
[2,]     1     4
[3,]     2     4
[4,]     1     3
[5,]     1     2
[6,]     3     4
```

```
isGYD(find.BIB(4, 6, 2))
```

```
[1] The design is a balanced incomplete block design w.r.t. rows.
```

```
lam = 3*1/3
lam
```

```
[1] 1
```

## R: BIBD example

- Make sure to put block first, or it's a wrong model. you should be sure to put the “treatment factor” after all “nuisance factors” in the aov model.

```
time <- c(73,74,71,75,67,72,73,75,68,75,72,75)
treatment <- c(rep(1,3),rep(2,3),rep(3,3),rep(4,3))
block <- c(1,2,4,2,3,4,1,2,3,1,3,4)
data.frame(treatment=treatment,block=block, time=time)
```

	treatment	block	time
1	1	1	73
2	1	2	74
3	1	4	71
4	2	2	75
5	2	3	67
6	2	4	72
7	3	1	73
8	3	2	75
9	3	3	68
10	4	1	75
11	4	3	72
12	4	4	75

```
# correct model
```

```
bibdmodel <- aov(time~factor(block)+factor(treatment))
summary(bibdmodel)
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
factor(block)	3	55.00	18.333	28.20	0.00147 **
factor(treatment)	3	22.75	7.583	11.67	0.01074 *
Residuals	5	3.25	0.650		

---  
Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

```
# incorrect model
```

```
bibdmodel_wrong <- aov(time~factor(treatment)+factor(block))
summary(bibdmodel_wrong)
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
factor(treatment)	3	11.67	3.889	5.983	0.041463 *
factor(block)	3	66.08	22.028	33.889	0.000953 ***
Residuals	5	3.25	0.650		

---  
Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

## Factorial Design with no blocking

- Key: More than one treatments and their interactions.
- In many scientific investigations, the interest lies in the study of effects of two or more factors simultaneously. Factorial designs are most commonly used for this type of investigation.
- Factorial design means that in each complete trial or replicate of the experiment all possible combinations of the levels of the factors are investigated. For example, if there are “a” levels of factor A and “b” levels of factor B, each replicate contains all “ab” treatment combinations.
- Main effect : the change in response produced by a change in the level of the factor. e.g.  $A = Y_{A+} - Y_{A-}$  and  $B = Y_{B+} - Y_{B-}$
- Interaction effect: the average difference in these two A effects or  $AB = ((Y_{A+B+} - Y_{A-B+}) + (Y_{A+B-} - Y_{A-B-}))/2$ . The interaction is large in this experiment which is also indicated by the non parallel lines. no interaction as indicated by the parallel lines.
- $y_{ijk} = \mu + \tau_i + \beta_j + (\tau\beta)_{ij} + \epsilon_{ijk}$  where  $i = 1 \dots a, j = 1 \dots b, k = 1 \dots n$ . Sum of  $\tau$  and  $\beta$  are 0, and so as the interaction terms w.r.t. i and j.
- $H_o : \tau_1 = \dots = \tau_a = 0$  and  $H_o : \beta_1 = \dots = \beta_b = 0$  and  $H_o : (\tau\beta)_{ij} = 0$
- Again ANOVA and F-test. GOOGLE IT. Three F test possible.

### R: Two factor factorial

- 2:70-1:70 and 3:70-1:70 most significant.
- some interaction in material type and temperature. Material 3 gives the best results.

```
material <- c(rep(1,12), rep(2,12), rep(3,12))
temp <- rep(c(rep(15,4), rep(70,4), rep(125,4)), 3)
response <- c(130,155,74,180,34,40,80,75,20,70,82,58,150,188,159,126,136,122,106,115,25,70,58,45,138,110)
model1 <- aov(response~factor(material)*factor(temp))
summary(model1)
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
factor(material)	2	10684	5342	7.911	0.00198 **
factor(temp)	2	39119	19559	28.968	1.91e-07 ***
factor(material):factor(temp)	4	9614	2403	3.560	0.01861 *
Residuals	27	18231	675		


---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

*# Model Adequacy*

```
par(mfrow=c(1,2))
plot(model1$residuals~model1$fitted.values)
qqnorm(model1$residuals);qqline(model1$residuals)
```





*# Multiple comparison, pick the one that's most effective.*

```
TukeyHSD(model1)
```

Tukey multiple comparisons of means  
95% family-wise confidence level

```
Fit: aov(formula = response ~ factor(material) * factor(temp))
```

```
$`factor(material)`
```

	diff	lwr	upr	p adj
2-1	25.16667	-1.135677	51.46901	0.0627571
3-1	41.91667	15.614323	68.21901	0.0014162
3-2	16.75000	-9.552344	43.05234	0.2717815

```
$`factor(temp)`
```

	diff	lwr	upr	p adj
70-15	-37.25000	-63.55234	-10.94766	0.0043788
125-15	-80.66667	-106.96901	-54.36432	0.0000001
125-70	-43.41667	-69.71901	-17.11432	0.0009787

```
$`factor(material):factor(temp)`
```


	diff	lwr	upr	p adj
2:15-1:15	21.00	-40.823184	82.823184	0.9616404
3:15-1:15	9.25	-52.573184	71.073184	0.9998527
1:70-1:15	-77.50	-139.323184	-15.676816	0.0065212
2:70-1:15	-15.00	-76.823184	46.823184	0.9953182
3:70-1:15	11.00	-50.823184	72.823184	0.9994703
1:125-1:15	-77.25	-139.073184	-15.426816	0.0067471
2:125-1:15	-85.25	-147.073184	-23.426816	0.0022351
3:125-1:15	-49.25	-111.073184	12.573184	0.2016535
3:15-2:15	-11.75	-73.573184	50.073184	0.9991463
1:70-2:15	-98.50	-160.323184	-36.676816	0.0003449
2:70-2:15	-36.00	-97.823184	25.823184	0.5819453
3:70-2:15	-10.00	-71.823184	51.823184	0.9997369
1:125-2:15	-98.25	-160.073184	-36.426816	0.0003574
2:125-2:15	-106.25	-168.073184	-44.426816	0.0001152
3:125-2:15	-70.25	-132.073184	-8.426816	0.0172076
1:70-3:15	-86.75	-148.573184	-24.926816	0.0018119
2:70-3:15	-24.25	-86.073184	37.573184	0.9165175
3:70-3:15	1.75	-60.073184	63.573184	1.0000000
1:125-3:15	-86.50	-148.323184	-24.676816	0.0018765

2:125-3:15	-94.50	-156.323184	-32.676816	0.0006078
3:125-3:15	-58.50	-120.323184	3.323184	0.0742711
2:70-1:70	62.50	0.676816	124.323184	0.0460388
3:70-1:70	88.50	26.676816	150.323184	0.0014173
1:125-1:70	0.25	-61.573184	62.073184	1.0000000
2:125-1:70	-7.75	-69.573184	54.073184	0.9999614
3:125-1:70	28.25	-33.573184	90.073184	0.8281938
3:70-2:70	26.00	-35.823184	87.823184	0.8822881
1:125-2:70	-62.25	-124.073184	-0.426816	0.0474675
2:125-2:70	-70.25	-132.073184	-8.426816	0.0172076
3:125-2:70	-34.25	-96.073184	27.573184	0.6420441
1:125-3:70	-88.25	-150.073184	-26.426816	0.0014679
2:125-3:70	-96.25	-158.073184	-34.426816	0.0004744
3:125-3:70	-60.25	-122.073184	1.573184	0.0604247
2:125-1:125	-8.00	-69.823184	53.823184	0.9999508
3:125-1:125	28.00	-33.823184	89.823184	0.8347331
3:125-2:125	36.00	-25.823184	97.823184	0.5819453

```
# interaction
```

```
interaction.plot(material, temp, response)
```

```
interaction.plot(temp, material, response)
```



```
par(mfrow=c(1,1))
```

## General Factorial Design with blocking

- Key: More than one treatments(A,B,C,..., where N=abc...) and their interactions, and blockings.
- ANOVA again for testing effect and interactions.
- For two-factor case with blocking,  $y_{ijk} = \mu + \tau_i + \beta_j + (\tau\beta)_{ij} + \delta_k + \epsilon_{ijk}$  where  $i = 1 \dots a$ ,  $j = 1 \dots b$ ,  $k = 1 \dots n$ . Sum of  $\tau$  and  $\beta$  are 0, and so as the interaction terms w.r.t. i and j.  $\delta$  is for blocking. If blocking has k levels, this is called  $2 \times k$  factorial design.
- If the treatments including their combinations and two blockings have the same levels, this is not only a factorial design, but also just pxp latin square. e.g., 2 filter and 3 cutter = 6 combinations, 6 days, 6 operators. In this case,  $y_{ijkl} = \mu + \alpha_i + \tau_j + \beta_k + (\tau\beta)_{jk} + \theta_l + \epsilon_{ijkl}$  where  $i = 1 \dots 6$ ,  $j = 1, 2, 3$ ,  $k = 1, 2$   $l = 1 \dots, 6$  and  $\tau, \beta$  are cutter and filter.  $\alpha$  and  $\theta$  are days and operators.

### R: three factor factorial

- As level increase, the average increase significantly. However, no interactions can be seen.

```
carb <- c(rep(10,8),rep(12,8),rep(14,8))
pressure <- rep(c(rep(25,4),rep(30,4)),3)
speed <- rep(c(200,200,250,250),6)
response <- c(-3,-1,-1,0,-1,0,1,1,0,1,2,1,2,3,6,5,5,4,7,6,7,9,10,11)
model <- aov(response~factor(carb)*factor(pressure)*factor(speed))
summary(model)
```

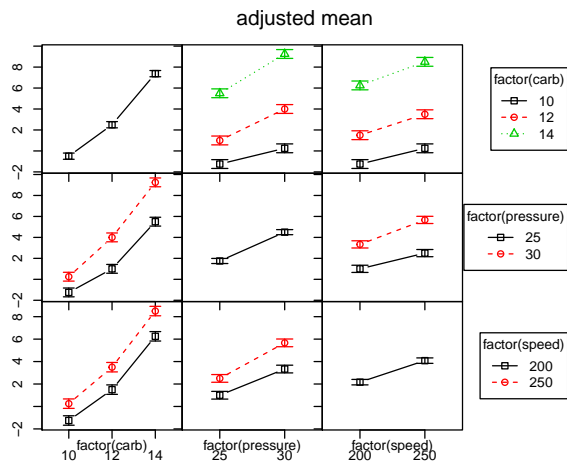
	Df	Sum Sq	Mean Sq	F value
factor(carb)	2	252.75	126.38	178.412
factor(pressure)	1	45.37	45.37	64.059
factor(speed)	1	22.04	22.04	31.118
factor(carb):factor(pressure)	2	5.25	2.63	3.706
factor(carb):factor(speed)	2	0.58	0.29	0.412
factor(pressure):factor(speed)	1	1.04	1.04	1.471
factor(carb):factor(pressure):factor(speed)	2	1.08	0.54	0.765
Residuals	12	8.50	0.71	

	Pr(>F)
factor(carb)	1.19e-09 ***
factor(pressure)	3.74e-06 ***
factor(speed)	0.00012 ***
factor(carb):factor(pressure)	0.05581 .
factor(carb):factor(speed)	0.67149
factor(pressure):factor(speed)	0.24859
factor(carb):factor(pressure):factor(speed)	0.48687
Residuals	

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

```
library(phia)
miplot <- interactionMeans(model)
plot(miplot)
```



## R: two factor with one blocking

- here, the order doesnt matter because of \*

```
Clutter <- rep(1:3,8)
Filter <- rep(c(1,1,1,2,2,2),4)
Operator <- c(rep(1,6),rep(2,6),rep(3,6),rep(4,6))
Response <- c(90,102,114,86,87,93,96,106,112,84,90,91,100,105,108,92,97,95,92,96,98,81,80,83)
model1 <- aov(Response~factor(Clutter)*factor(Filter)+factor(Operator))
summary(model1)
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)	
factor(Clutter)	2	335.6	167.8	15.132	0.000253	***
factor(Filter)	1	1066.7	1066.7	96.192	6.45e-08	***
factor(Operator)	3	402.2	134.1	12.089	0.000277	***
factor(Clutter):factor(Filter)	2	77.1	38.5	3.476	0.057507	.
Residuals	15	166.3	11.1			

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

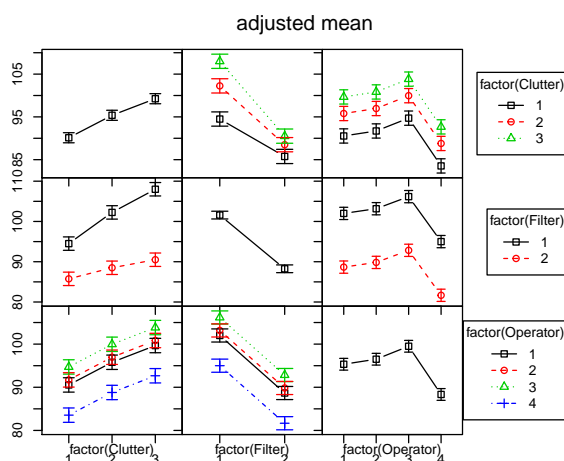
```
model2 <- aov(Response~factor(Operator)+factor(Clutter)*factor(Filter))
summary(model2)
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)	
factor(Operator)	3	402.2	134.1	12.089	0.000277	***
factor(Clutter)	2	335.6	167.8	15.132	0.000253	***
factor(Filter)	1	1066.7	1066.7	96.192	6.45e-08	***
factor(Clutter):factor(Filter)	2	77.1	38.5	3.476	0.057507	.
Residuals	15	166.3	11.1			

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

```
miplot1 <- interactionMeans(model1)
plot(miplot1)
```



## $2^k$ Factorial Design

- Key:  $k$  factors, each at only two levels.
- Assumptions: factors are fixed, designed are completely randomized, the usual normality assumptions.
- Particularly useful in early stages of experiments when many factors are likely to be investigated: It provides the smallest number of runs with which  $k$  factors can be studied in a complete factorial design. Since there are only two levels for each factor, we assume that the response is approximately linear over the range of the factor levels chosen.
- $N = 2^k \times r$ , where  $r$  is the number of replicates for each treatment factors.
- For  $2^2$  design, we consider main effects. e.g.  $A = \frac{1}{2}(\frac{a-(1)}{r} + \frac{ab-b}{r})$ ,  $AB = \frac{1}{2}(\frac{ab-b}{r} - \frac{a-(1)}{r})$ . If they are positive/negative, increase/decrease in means are yielded. if  $AB$  is small, small interaction. Use ANOVA to see those main effects are significant.
- $2^k$  factorial design can be expressed in regression model.  $y = \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \epsilon$
- For  $2^3$  factorial design,  $A = \frac{1}{4r}(a-(1)+ab-b+ac-c+abc-bc)$ ,  $AB = \frac{1}{4r}(abc+ab+c+(1)-ac-bc-a-b)$ ,  $ABC = \frac{1}{4r}(abc+a+b+c-ab-ac-bc-(1))$ .  $SS = \frac{Contrast^2}{8r}$ , where  $Contrast_{AB...K} = (a \pm 1)(b \pm 1) \dots (k \pm 1)$  e.g.  $Contrast_{AB} = (a-1)(b-1)(c+1)$ .
- $AB \dots K = \frac{2}{r2^k}(Contrast_{AB...K})$ ,  $SS_{AB...K} = \frac{1}{r2^k}(Contrast_{AB...K})^2$
- General  $2^k$  factorial design has:  $\binom{k}{2}$  two factor interactions,  $\binom{k}{3}$  three factor interactions and so forth. Procedure:
  1. Estimate factor effects and examine their signs and magnitudes : This gives the experimenter preliminary information regarding which factors and interactions may be important and in which directions these factors should be adjusted to improve the response.
  2. Form initial model: (a) If the design is replicated, fit the full model. (b) If there is no replication, form the model using a normal probability plot of the effects.
  3. Perform statistical testing : use the analysis of variance to formally test for the significance of main effects and interaction.
  4. Refine model : remove any nonsignificant variables from the full model.
  5. Analyze residuals : check for model adequacy and assumptions.
  6. Interpret results : include graphical analysis main effect or interaction plots.
- Some problems that arise with a single replicate ( $r=1$ ): We may be fitting a model to noise. If the response is highly variable, misleading conclusions may result from the experiment. Fitting the model results in zero degrees of freedom for error.
- That is, we cannot use ANOVA to test this. In stead, examine a Normal Probability plot of the estimates of the effects. All of the effects that lie along normal QQ plot line is negligible, whereas the large effects are far from the line.

### R: $2^4$ factorial design

```
A <- rep(c(-1,1),8)
B <- rep(c(-1,-1,1,1),4)
C <- rep(c(rep(-1,4),rep(1,4)),2)
D <- c(rep(-1,8),rep(1,8))
response <- c(45,71,48,65,68,60,80,65,43,100,45,104,75,86,70,96)
model.sg <- aov(response~factor(A)*factor(B)*factor(C)*factor(D))
summary(model.sg)
```

	Df	Sum Sq	Mean Sq
factor(A)	1	1870.6	1870.6
factor(B)	1	39.1	39.1
factor(C)	1	390.1	390.1
factor(D)	1	855.6	855.6
factor(A):factor(B)	1	0.1	0.1
factor(A):factor(C)	1	1314.1	1314.1
factor(B):factor(C)	1	22.6	22.6
factor(A):factor(D)	1	1105.6	1105.6
factor(B):factor(D)	1	0.6	0.6
factor(C):factor(D)	1	5.1	5.1
factor(A):factor(B):factor(C)	1	14.1	14.1
factor(A):factor(B):factor(D)	1	68.1	68.1
factor(A):factor(C):factor(D)	1	10.6	10.6
factor(B):factor(C):factor(D)	1	27.6	27.6
factor(A):factor(B):factor(C):factor(D)	1	7.6	7.6


---

### Normal plot for effects

```
model.lm <- lm(response~A*B*C*D)
effect <- 2*(coef(model.lm)[-1])
effect
```

A	B	C	D	A:B	A:C	B:C	A:D	B:D
21.625	3.125	9.875	14.625	0.125	-18.125	2.375	16.625	-0.375
C:D	A:B:C	A:B:D	A:C:D	B:C:D	A:B:C:D			
-1.125	1.875	4.125	-1.625	-2.625	1.375			

```
qqnorm(effect)
qqline(effect)
```




### Main/Interaction plots

```
library(gplots)
par(mfrow=c(1,3))
```

```

plotmeans(response~A,xlab="Factor A",ylab="Average Response", p=0, main="Main effect Plot",barcol="black",
plotmeans(response~C,xlab="Factor A",ylab="Average Response", p=0, main="Main effect Plot",barcol="black",
plotmeans(response~D,xlab="Factor A",ylab="Average Response", p=0, main="Main effect Plot",barcol="black",

```



```

par(mfrow=c(1,1))

```



## 2<sup>k</sup> Factorial Design with blockings (confoundings)

- Key: k factors, each at only two levels and multiple blockings each two levels.
- Look at the signs for each components of factorial main effects. Here, AB is confounded with blocks because the its signs and the number of blocks matches ( + for 1, - for 2). Typically confound the highest order interaction with blocks.

Treatments	I	A	B	AB	Block
(1)	+	-	-	+	1
a	+	+	-	-	2
b	+	-	+	-	2
ab	+	+	+	+	1

### R: 2<sup>4</sup> with no blocking

```
A <- rep(c(-1,1),8)
B <- rep(c(-1,-1,1,1),4)
C <- rep(c(rep(-1,4),rep(1,4)),2)
D <- c(rep(-1,8),rep(1,8))
response <- c(45,71,48,65,68,60,80,65,43,100,45,104,75,86,70,96)
modell1 <- aov(response~factor(A)*factor(B)*factor(C)*factor(D))
summary(modell1)
```

```
              Df Sum Sq Mean Sq
factor(A)      1 1870.6   1870.6
factor(B)      1   39.1     39.1
factor(C)      1  390.1    390.1
factor(D)      1  855.6    855.6
factor(A):factor(B) 1    0.1      0.1
factor(A):factor(C) 1 1314.1   1314.1
factor(B):factor(C) 1   22.6     22.6
factor(A):factor(D) 1 1105.6   1105.6
factor(B):factor(D) 1    0.6      0.6
factor(C):factor(D) 1    5.1      5.1
factor(A):factor(B):factor(C) 1   14.1    14.1
factor(A):factor(B):factor(D) 1   68.1    68.1
factor(A):factor(C):factor(D) 1   10.6    10.6
factor(B):factor(C):factor(D) 1   27.6    27.6
factor(A):factor(B):factor(C):factor(D) 1    7.6     7.6
```

```
# effect estimated
```

```
effect1 <- 2*coef(lm(response~A*B*C*D))[-1]
effect1
```

```
      A      B      C      D      A:B      A:C      B:C      A:D      B:D
21.625  3.125  9.875 14.625  0.125 -18.125  2.375 16.625 -0.375
  C:D  A:B:C  A:B:D  A:C:D  B:C:D A:B:C:D
-1.125  1.875  4.125 -1.625 -2.625  1.375
```

### R: 2<sup>4</sup> with blocking

- ANOVA does not print the ABCD interaction, since it is confounded with blocking. Also, the estimated block effect includes ABCD interaction effect.

```
# modification 1 : add a block #
block <- A*B*C*D

# modification 2 : responses in block 1 has -20 lower units than original
response[block==1] <- response[block==1]-20 #modified

model2 <- aov(response~factor(A)*factor(B)*factor(C)*factor(D)+factor(block))
summary(model2)
```

	Df	Sum Sq	Mean Sq
factor(A)	1	1870.6	1870.6
factor(B)	1	39.1	39.1
factor(C)	1	390.1	390.1
factor(D)	1	855.6	855.6
factor(block)	1	1387.6	1387.6
factor(A):factor(B)	1	0.1	0.1
factor(A):factor(C)	1	1314.1	1314.1
factor(B):factor(C)	1	22.6	22.6
factor(A):factor(D)	1	1105.6	1105.6
factor(B):factor(D)	1	0.6	0.6
factor(C):factor(D)	1	5.1	5.1
factor(A):factor(B):factor(C)	1	14.1	14.1
factor(A):factor(B):factor(D)	1	68.1	68.1
factor(A):factor(C):factor(D)	1	10.6	10.6
factor(B):factor(C):factor(D)	1	27.6	27.6

```
# Estimated effects
effect2 <- 2*coef(lm(response~A*B*C*D+block))[-1]
effect2
```

A	B	C	D	block	A:B	A:C	B:C	A:D
21.625	3.125	9.875	14.625	-18.625	0.125	-18.125	2.375	16.625
B:D	C:D	A:B:C	A:B:D	A:C:D	B:C:D	A:B:C:D		
-0.375	-1.125	1.875	4.125	-1.625	-2.625	NA		

```
# ANOVA for the final model (reduced model including only significant factors)
```

```
model3 <- aov(response~factor(A)+factor(C)+factor(D)+factor(A)*factor(C)+factor(A)*factor(D)+factor(blo
summary(model3)
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
factor(A)	1	1870.6	1870.6	89.76	5.60e-06 ***
factor(C)	1	390.1	390.1	18.72	0.001915 **
factor(D)	1	855.6	855.6	41.05	0.000124 ***
factor(block)	1	1387.6	1387.6	66.58	1.89e-05 ***
factor(A):factor(C)	1	1314.1	1314.1	63.05	2.35e-05 ***
factor(A):factor(D)	1	1105.6	1105.6	53.05	4.65e-05 ***
Residuals	9	187.6	20.8		

```
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

## $2^{\{k-p\}}$ Fractional Factorial Design

- Key:  $k$  factors, each at only two levels, but too many  $K$ 's.
- As the number of factors in a  $2^k$  factorial design increases, the number of runs required for a complete replicate of the design rapidly outgrows the resources of most experimenters. Thus use fractional factorial design. Three properties:
  1. The sparsity of effects principle : When there are several variables, the process is likely to be driven primarily by some of the main effects and low order interactions.
  2. The projection property : Fractional factorial designs can be projected into stronger (larger ) designs in the subset of significant factors.
  3. Sequential experimentation : It is possible to combine the runs of two (or more) fractional factorials to resolve difficulties in interpretation.
- The highest interaction that has the identical effect sign as the identity  $I$  is called the generator, and for instance, if the signs for  $ABC$  is all plus and so as  $I$ , then  $I = ABC$  is called the defining relation for our design.
- Then, we have;  $[A] = \frac{1}{2}(a - b - c + abc)$ ,  $[B] = \frac{1}{2}(-a + b - c + abc)$ ,  $[C] = \frac{1}{2}(-a - b + c + abc)$ . Likewise,  $[BC] = \frac{1}{2}(a - b - c + abc)$ ,  $[AC] = \frac{1}{2}(-a + b - c + abc)$ ,  $[AB] = \frac{1}{2}(-a - b + c + abc)$ . Thus,  $[A] = [BC]$ ,  $[B] = [AC]$ ,  $[C] = [AB]$ . It is impossible to differentiate those.
- In fact,  $[A] \rightarrow A + BC$ ,  $[B] \rightarrow B + AC$ ,  $[C] \rightarrow C + AB$
- $A \cdot I = A \cdot ABC = A^2BC \rightarrow A = BC$
- If we take the “other half”,  $[A]' \rightarrow A - BC$ ,  $[B]' \rightarrow B - AC$ ,  $[C]' \rightarrow C - AB$
- An important property of a fractional design is its resolution or ability to separate main effects and low order interactions from one another. The resolution of the design is the minimum word length in the defining relation excluding (1). e.g. if  $I=ABC$ , the resolution is 3.
- Resolution 3: Good for initial screening. The higher the resolution, the less restrictive the assumptions that are required regarding which interactions are negligible to obtain a unique interpretation of the results.

# Computational Techniques

## Standarization / Normalization

## K-Fold Cross Validation

- A method to validate the model and its parameters. MSE is highly variable measurement for the model. Thus, cross validation reduce this variability and help us obtain the most stable MSE.
  - Devide the data set into K different parts. Remove the first part, fit a model using the rest of the parts, and test on this removed first part and take MSE. Repeat this procedure for all folds. At last, average all K different MSE and obtain the MSE for the model.
  - When  $K = n$ , it's called "Leace-One-Out Cross Validation".
- 

### R: CV to choose order of polynomial

```
library(boot)
set.seed(1)
K = 5
CV_MSE = rep(0,K)
for(i in 1:K){
  glm.fit = glm(Sepal.Width~poly(Sepal.Length,i), data=iris)
  CV_MSE[i] = cv.glm(iris, glm.fit)$delta[1]
}
CV_MSE
```

```
[1] 0.1908667 0.1890806 0.1859583 0.1863217 0.1926010
```

**R: Package “crossval”**

```
library(crossval)
```

## Bootstrap

**R: Obtain mean and CI using package “boot”**

```
library(boot)
set.seed(1)

# Mean
mean.fn = function(data, index){
  return(c(mean(data[index,1])))
}
mean_boot = boot(iris, mean.fn, R=1000)

# CI
CI_norm_boot = boot.ci(mean_boot, type = 'norm')
CI_norm_boot$normal

      conf
[1,] 0.95 5.708369 5.976465
```

# R Features

## Basics

## Functions



## I() and poly() for polynomial regression

- I() function simply grants lm() function to add higher order variables.
- poly() function uses an orthogonal basis to fit polynomial regression. That is, the variables are linearly transformed into linearly independent set, thus no multicollinearity is present. Moreover, the statistics generated such as  $r^2$  and resulting plots are the same as using I() functions. However, the coefficients generated are different hence it suffers from interpretations.

```

grids=seq(from=min(iris$Sepal.Length),to=max(iris$Sepal.Length), by=0.01)

m_I = lm(Sepal.Width~Sepal.Length+I(Sepal.Length^2), data=iris)
summary(m_I)

```

Call:

```
lm(formula = Sepal.Width ~ Sepal.Length + I(Sepal.Length^2),
    data = iris)
```

Residuals:

Min	1Q	Median	3Q	Max
-1.13070	-0.26310	-0.02446	0.25728	1.38725

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )	
(Intercept)	6.41584	1.58499	4.048	8.33e-05	***
Sepal.Length	-1.08556	0.53625	-2.024	0.0447	*
I(Sepal.Length^2)	0.08571	0.04476	1.915	0.0574	.

---  
 Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4304 on 147 degrees of freedom


Multiple R-squared: 0.03783, Adjusted R-squared: 0.02474

F-statistic: 2.89 on 2 and 147 DF, p-value: 0.05877

```

pred = predict(m_I,newdata=list(Sepal.Length=grids),se=TRUE)
plot(Sepal.Width~Sepal.Length, data=iris);lines(grids,pred$fit, lwd=2, col="blue")

```



```

m_poly = lm(Sepal.Width~poly(Sepal.Length,2), data=iris)
summary(m_poly)

```

Call:

```
lm(formula = Sepal.Width ~ poly(Sepal.Length, 2), data = iris)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-1.13070	-0.26310	-0.02446	0.25728	1.38725

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	3.05733	0.03515	86.991	<2e-16 ***
poly(Sepal.Length, 2)1	-0.62552	0.43044	-1.453	0.1483
poly(Sepal.Length, 2)2	0.82430	0.43044	1.915	0.0574 .

---


Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.4304 on 147 degrees of freedom

Multiple R-squared: 0.03783, Adjusted R-squared: 0.02474

F-statistic: 2.89 on 2 and 147 DF, p-value: 0.05877

```
pred = predict(m_poly, newdata=list(Sepal.Length=grids), se=TRUE)
plot(Sepal.Width~Sepal.Length, data=iris); lines(grids, pred$fit, lwd=2, col="blue")
```



## Normalization and Standarization

`apply()`, `sapply()`, `tapply()`, `lapply()`

Package: `dplyr`